User Guide & primer

A plain-language guide to using this pack **and** to the ideas behind it. It assumes no prior knowledge of business architecture. If the words *value stream*, *capability* or *data product* are new to you, start at “The big idea” and read straight down — by the end you'll be able to look at any AI initiative and say exactly how it connects to a business outcome, what it depends on, and how you'd know it worked. That is the whole point of the pack: to make Data & AI decisions legible to the people who sponsor them.

1 · The big idea — why AI needs a business architecture

Most AI initiatives fail for a business reason, not a technical one.

They get built without a clear line back to a decision that matters, an outcome someone owns, or data anyone trusts. A **business architecture** is simply the map that draws those lines. It lets you point at any AI idea and ask: *whom does it serve, what decision does it change, what result does it move, and what does it depend on?* If the map has a gap, that is where the initiative will struggle.

This pack lays out one such map for AGGPSA's ecosystem work. Everything in it hangs off a single chain we call the **golden thread** — a line from a person the ecosystem serves, all the way down to the governed data an AI system runs on:

Stakeholder → Value Proposition → Business Outcome + KPI → Value Stream → Value Stage → Customer Journey → Capability → Process → Decision → AI Use-Case → AI Agent → Data Product → Data Domain

WHAT YOU CAN DO WITH THE THREAD

- **Justify** an AI initiative: trace it up to the outcome and KPI it moves.
- **De-risk** it: trace it down to the data it needs and the human oversight it keeps.
- **Prioritise**: compare initiatives on the value they create versus the risk they carry.
- **Align** a room: give everyone the same words for the same things.

HOW TO READ THIS GUIDE

- Section 2 defines every building block in plain language, with a running example.
- Section 3 is a short playbook for taking one AI idea through the model.
- Then: how to read the charts, how to move around, and what each tab is for.
- No jargon is assumed — terms are defined the first time they appear.

> Prefer to learn by doing? Open the Navigator and walk one thread end-to-end. Pick a stakeholder and keep clicking forward — the guide below explains what each step means as you go.

2 · The building blocks, in plain language

These are the thirteen links of the golden thread, in order. For each one: what it is, the question it answers, why it matters to a sponsor, and a worked example from a single running story — **improving entrepreneur survival & scale**. Read them top to bottom and the whole model falls into place. Wherever AI enters, it is called out in amber.

1

Stakeholder

Who is this for?

What it is. The people or groups the ecosystem work serves or must answer to — beneficiaries and entrepreneurs, ESOs and partners, funders and regulators, and internal teams like Programmes or Data & Research.

Why it matters. An AI initiative that can't name the stakeholder it helps is a solution looking for a problem. Everything in the model traces back to a stakeholder, which is why the walk starts here.

Example. *Beneficiaries & entrepreneurs — the people AGGPSA's ecosystem work exists to support and must answer to.*

2

Value Proposition

What do we promise them?

What it is. The value proposition — the promise of value made to that stakeholder. What they get, and why it is worth their attention.

Why it matters. It anchors every downstream cost and project to a benefit someone actually wants. No promise, no reason to build.

Example. *“Right-fit catalytic support at the right moment, through the right intermediary.”*

3

Business Outcome

What result are we trying to move — and how will we know?

What it is. The business outcome is the measurable result that proves the promise is being kept. Each outcome carries KPIs — the specific numbers that quantify it.

Why it matters. This is where AI has to earn its place. If an initiative can't name the outcome and KPI it moves, it can't be justified or measured. Outcomes and KPIs are the language executives already speak.

Example. *Outcome: improve entrepreneur survival & scale. KPI: survival rate, and jobs created.*

4

Value Stream

How does the business create this value, start to finish?

What it is. A value stream is the end-to-end set of activities that creates and delivers the value — from the trigger to the result — seen from the organisation's side.

Why it matters. It gives an AI investment somewhere concrete to live. You don't improve “the business” in the abstract; you improve a stream.

Example. *Support Entrepreneurs & ESOs — everything from first contact to a resilient, growing enterprise.*

5

Value Stage

Which specific step are we improving?

What it is. A value stage is one step of the stream, with a clear start and finish and its own KPI. Stages are how a long stream is broken into measurable pieces.

Why it matters. AI is applied stage by stage. Because each stage's KPI rolls up into the outcome, you can see exactly where value is won or lost — and target the weakest step.

Example. *Sustain — the stage where enterprises are kept alive and growing. KPI: survival rate.*

6

Customer Journey

What does this feel like for the customer?

What it is. The beneficiary journey is the same value seen from the beneficiary's side — the experience they actually go through, moment by moment.

Why it matters. Analytics that improves an internal number but worsens the experience destroys value. Modelling the journey alongside the stream keeps the beneficiary in the picture.

Example. *A township entrepreneur struggles quietly; an ESO notices and connects timely, relevant support — rather than a generic programme.*

7

Capability

What must we be able to do?

What it is. A capability is an ability the business must have to deliver — a “what we can do,” independent of how teams are organised or which tools they use.

Why it matters. Capabilities are reusable building blocks. One AI investment in a capability can serve many streams and journeys, which is how you avoid re-solving the same problem five times.

Example. *Enterprise Support and Ecosystem Intelligence — the abilities to understand, target and grow enterprise support.*

8

Business Process

How exactly do we do it today?

What it is. A business process is the concrete, repeatable way a capability actually gets done — its steps, inputs and outputs. The pack shows each as a SIPOC (Suppliers, Inputs, Process, Outputs, Customers).

Why it matters. You cannot automate what you cannot describe. The process is the real work where AI plugs in — and where you see what a model would change.

Example. *Support Targeting — the process that scores enterprises for need and potential, then triggers the right support.*

9

Decision

What is being decided, and on what basis?

What it is. A decision is a single judgement made inside a process: inputs go in, an outcome comes out, following rules. The pack shows each as a decision table (if these conditions, then this outcome).

Why it matters. Most AI automates or assists a decision. Naming the decision is the exact moment a vague “AI initiative” becomes something concrete you can build, test and govern.

Example. *Support Decision — is this enterprise viable and underserved, and what is the best support for them?*

> **Where AI comes in.** This is the hand-off point to AI: a rule-based decision becomes a candidate for a model that decides faster, at scale, and more accurately.

10

AI Use Case

What is the AI actually doing for us?

What it is. An AI use-case is a specific application of AI to a decision or task — with a value it creates and a risk it carries. This is the “AI initiative” itself.

Why it matters. By this point the initiative is anchored to a decision, a stage, an outcome and a KPI — so its value and its risk are both visible and comparable against every other use-case, instead of being sold on hype.

Example. *Enterprise survival prediction & support — predict who is at risk and recommend the best support.*

> **Where AI comes in.** In this pack each use-case is scored 1–5 on value (growth, relevance, retention, CX, cost) and on risk (data readiness, model, integration, regulatory) — so a portfolio can be prioritised on evidence.

11

AI Agent

What runs it, and who oversees it?

What it is. An AI agent is the working system that runs a use-case — the model or service that does the job — together with its stance on human oversight.

Why it matters. Executives need to know not just that AI is used, but how it is operated and controlled. The agent is where accountability for the running system lives.

Example. *Diagnostic & Targeting agent — scores enterprises; high-value support is routed to a programme officer for approval.*

> **Where AI comes in.** Human-in-the-Loop: for high-stakes or high-risk outputs, a person reviews and approves before anything acts. It is the main dial between full automation and human control.

12

Data Product

What data does this depend on — and can we trust it?

What it is. A data product is a governed, owned, reusable data asset the AI needs, with a clear owner, an agreed definition, and a quality/contract. Not a raw table — a managed product.

Why it matters. AI is only as good as the data beneath it. Data products make that dependency explicit and accountable — often the real difference between a demo that impresses and a system that survives in production.

Example. *Township Enterprise Data and Impact Model Store — the trusted data the targeting model consumes and the scores it produces.*

13

Data Domain

Who is accountable for this data?

What it is. A data domain is the area of the business that owns and governs a group of related data products — with a named owner and steward.

Why it matters. It puts a name against the data. Ownership and stewardship are what make data trustworthy at scale, and what a regulator or auditor will ask for first.

Example. *Impact & Evidence — the domain that owns enterprise scores and model outputs across programmes.*

A few supporting terms

Words you'll meet across the pack that aren't their own step in the walk.

PERSONA

A named, representative example of a stakeholder — e.g. a specific beneficiary segment. Personas make an abstract stakeholder concrete so teams design for a real person, not an average.

SEMANTIC MODEL / GOVERNED TERM

The agreed, written definition of a business term — what “survival”, “active enterprise” or “catalytic leverage” actually mean. When everyone's analytics and every report uses the same definition, the numbers reconcile. This is the vocabulary layer beneath the data products.

ECOSYSTEM TARGETING ENGINE

The engine that unifies ecosystem and beneficiary signals in near-real-time and activates them — the plumbing that lets the right support reach the right beneficiary at the right moment. In the model it is a set of capabilities and services the targeting use-cases draw on.

HUMAN-IN-THE-LOOP (HITL)

A control where a person reviews or approves an analytics or AI output before it takes effect. Used for high-value, sensitive or regulated decisions. The presence or absence of HITL is one of the clearest signals of how much trust you are placing in a model.

The one sentence to remember: a good AI initiative is a *use-case* that automates a *decision* inside a *process*, which delivers a *capability* a *value stage* needs, moving a *KPI* for an *outcome* promised to a *stakeholder* — running on a governed *data product*, with a human in the loop where the stakes are high.

3 · How to apply AI to a business problem

A simple, repeatable way to take one idea through the model. Do it on paper, or do it live in the Navigator. The example in blue continues the survival story.

1

Start from an outcome, not a technology

Name the systemic result you want to move and the KPI that measures it. Resist starting with “let's use GenAI.” Start with the number.

Outcome: improve entrepreneur survival & scale · KPI: survival rate.

2

Find the decision that moves it

Trace to the value stage that most affects that KPI, then to the decision made there. Analytics acts on decisions — so this is where a real initiative forms.

Stage: Sustain → Decision: which enterprises are at risk and what is the best support?

3

Name the analytics use-case

State plainly what the analytics will do for that decision, and score it: how much value (systemic, evidence, leverage, reach, efficiency) against how much risk (data, model, equity, safeguarding, adoption).

Use-case: enterprise survival prediction & support · high value, moderate risk.

4

Check the data it depends on

List the data products the use-case needs and whether they are owned, defined and trustworthy. This is where most initiatives quietly fail — surface it early.

Needs: Township Enterprise Data, model scores · owner: Impact & Evidence domain.

5

Decide the human oversight

Choose where a person stays in the loop. High-value or sensitive actions get human approval; low-stakes, high-volume ones can run automatically.

Auto-score everyone; a programme officer approves support above a value threshold.

6

Close the loop on the KPI

Ship it, then watch the same KPI you started with. If it moves, you have proof; if not, you have a precise place to look. The thread makes the feedback loop obvious.

Track survival on the supported group vs a control; feed results back into targeting.

> Notice what the playbook prevents: an AI project with no owned outcome, no named decision, or no trustworthy data. Those three gaps account for most failed initiatives — and the model surfaces all three before a line of code is written.

4 · Reading the visualisations

The pack shows the same model several ways. Here is how to read each one.

NAVIGATOR — THE GUIDED WALK

- Pick an element at each step; the next step shows only what connects to it.
- The numbered stepper on the left is your position along the thread.
- Best for learning the model and answering “what connects to what.”

GRAPH — HIERARCHY TREE

- Expand any node (▶) to drill down a level at a time.
- Click a node for its plain-language explanation and its value.
- Best for exploring freely rather than following one path.

GRAPH — VALUE FLOW (SANKEY)

- Ribbons flow left → right from Stakeholder down to Data Product.
- **Ribbon width = how much value flows** (AI value score, or KPI weight — your choice).
- Wider bands are where value concentrates. Use the filter rail to hold “all” or a few items at any level.

AI USE-CASES — VALUE VS RISK

- Each bubble is a use-case; up-and-right = high value.
- Watch the risk axis: a high-value, high-risk idea needs stronger controls.
- Best for prioritising a portfolio on evidence, not opinion.

ARCHITECTURE — CAPABILITY HEATMAP

- Tiles are capabilities; colour shows attention needed or maturity.
- Hot tiles are where investment (including AI) will pay back most.

VALUE STREAMS — GOVERNANCE VIEW

- Traces a stream down to who owns and governs it.
- Best for answering “who is accountable” for an outcome or control.

5 · Finding your way around

THE LEFT SIDEBAR

- Every page is a numbered tab down the left; the current one is highlighted.
- The Modelware brand, the data *Source* selector and the version stamp live in the rail.
- On a narrow screen the rail folds into a bar across the top.

THE DATA SOURCE SELECTOR

- Switches which data the whole pack shows.
- *Generic (baked-in)* is the published model.
- *Imported (Excel)* appears after you import a workbook on Model I/O.

THE TRACEABILITY PATH

- Shown on the **Navigator** — it is the thread you are building.
- It fills in as you choose an element at each step.
- Other pages stay clean; your selection is still remembered behind the scenes.

CARRIED SELECTION

- A choice made in the Navigator follows you to other pages and focuses them.
- The Graph's Value Flow starts from it, and you can widen from there.
- Clear it with the Navigator's reset or the Graph's “Whole model”.

6 · Every page — purpose, architecture message & what you'll see

Every page in the app is documented below in navigation order — a live **screenshot**, what the page is **for**, the **architecture message** it presents (which layer of the model it shows and the point it makes), and what you'll see and do on it. Screenshots are from the **AGGPSA** source; the pack is source-switchable, so every layout is identical for other sources with that source's own content. Drop these straight into training slides.

Landing `index.html`

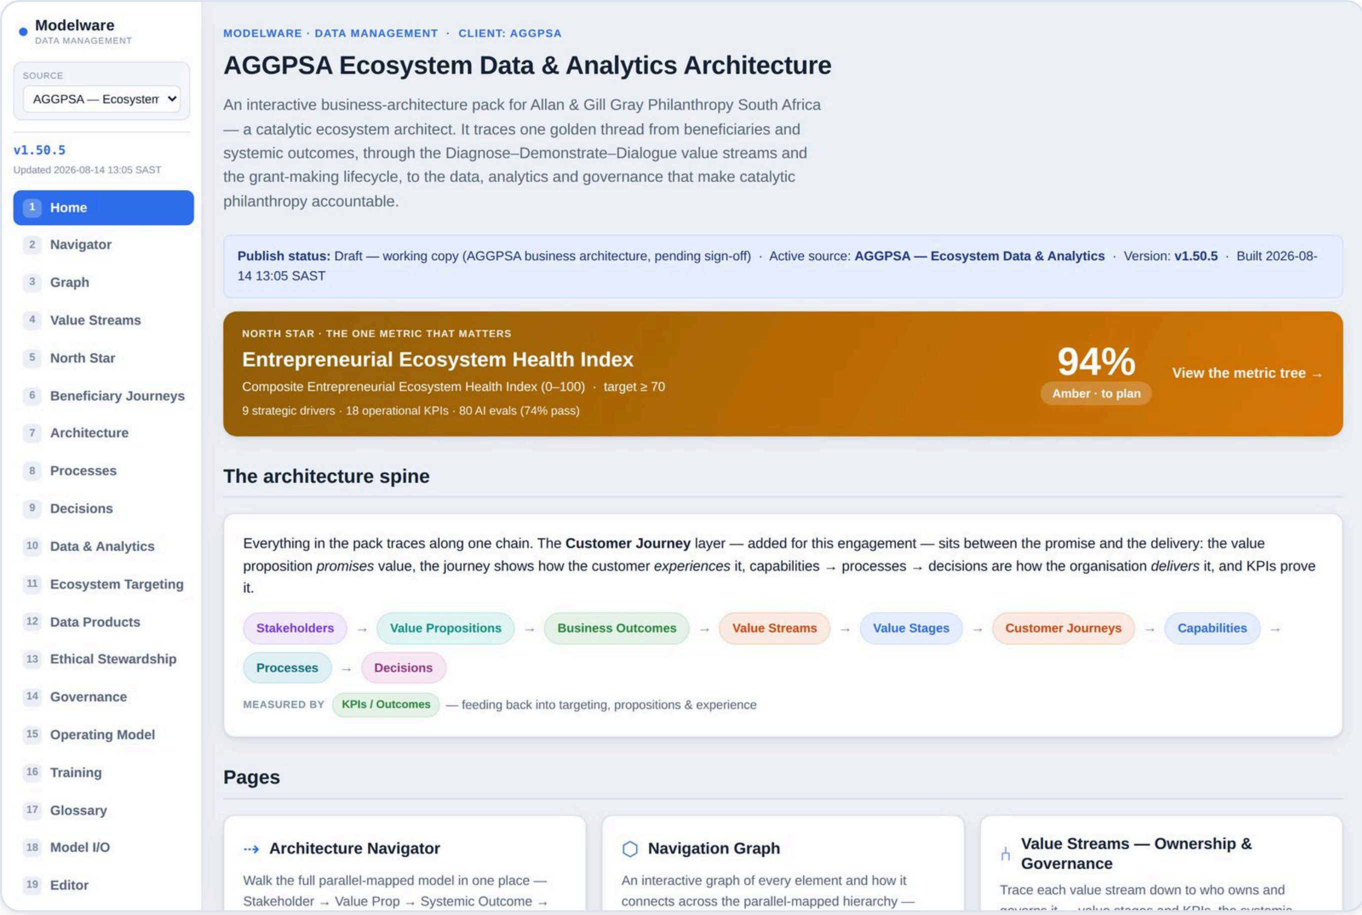

The front door. Explains the engagement, shows the architecture spine, lists every page with who it is for, links the downloadable workbooks and shows the current published version.

> **ARCHITECTURE MESSAGE**
>
> Orientation. Presents the whole architecture as one golden thread — Stakeholder → Value Proposition → Outcome & KPI → Value Stream → Value Stage → Journey → Capability → Process → Decision → AI Use-Case → AI Agent → Data Product → Domain — and frames every other page as one lens on that single connected model.

WHAT YOU'LL SEE

- The one-line architecture spine (Stakeholder → … → KPIs)
- A directory card for every page, tagged by audience
- Publish status, active source and version history

HOW TO USE IT

- Start here to get your bearings
- Click any directory card to open that page
- Check the version history to see what changed and when

Open Home →

Navigator

Architecture Navigator `architecture_navigator.html`

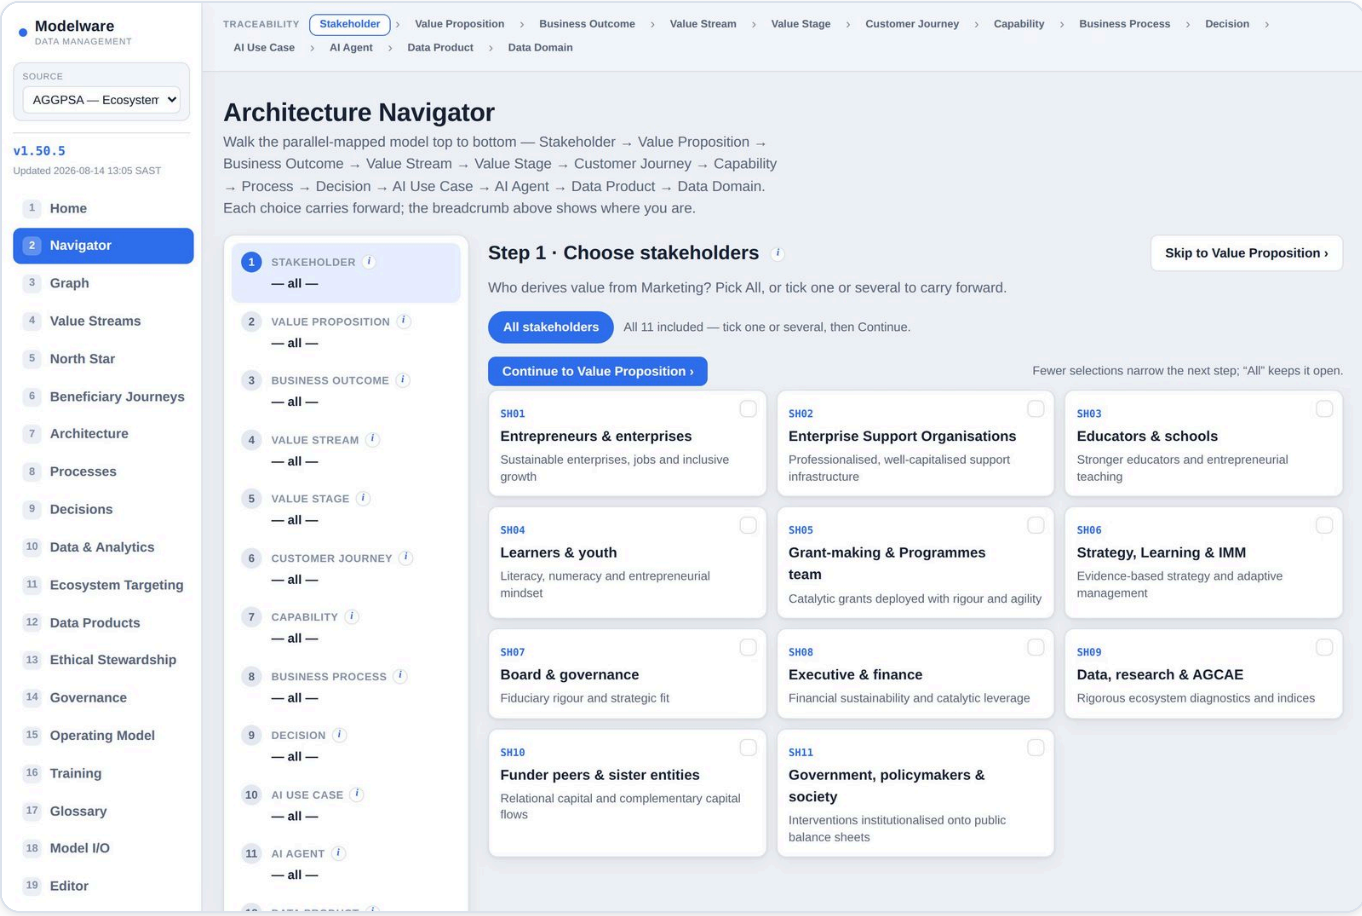

A guided, step-by-step walk through the whole parallel-mapped model. You pick a starting point and move one link at a time along the traceability chain, carrying your selection with you.

> **ARCHITECTURE MESSAGE**
>
> The traceability spine itself. Presents the parallel-mapped model as one navigable chain and proves that every element is connected end-to-end — from the person served to the governed data an AI runs on — with your selection carried the whole way.

WHAT YOU'LL SEE

- A stepper for the 13 levels of the chain
- At each step, the elements available given what you've already chosen
- A breadcrumb that fills in as you go

HOW TO USE IT

- Pick a Stakeholder (or any step) to begin
- Move forward and back along the chain; your selection is carried to every other page
- Use it to answer 'what connects to what' end-to-end

Open Navigator →

Graph

Navigation Graph `navigation_graph.html`

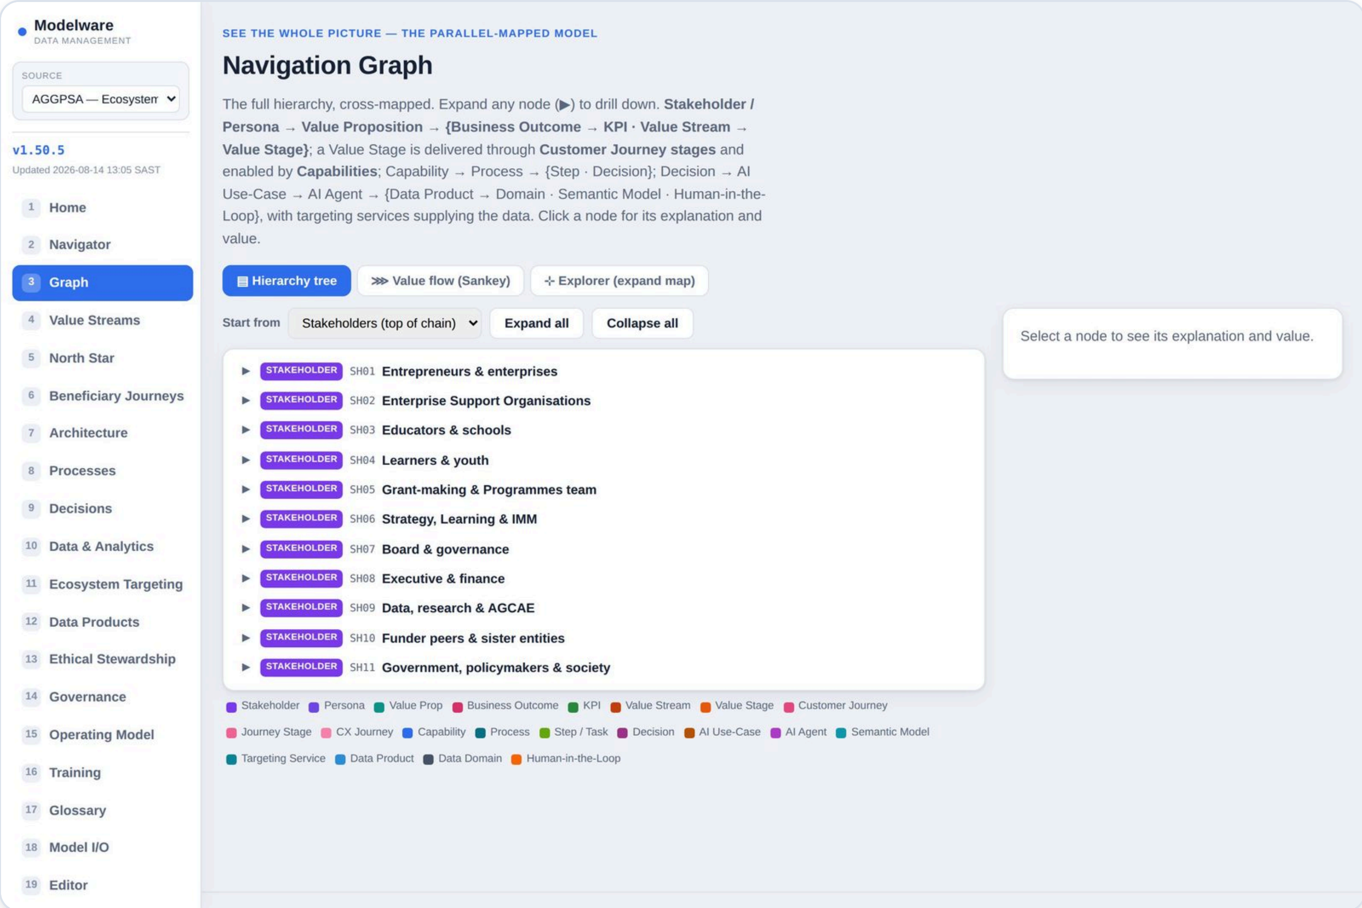

An interactive picture of the model. Two views: an expandable hierarchy tree, and a Value Flow (Sankey) that shows how value flows from stakeholders down to data products.

> **ARCHITECTURE MESSAGE**
>
> The model as a connected network and a value flow. Presents the same architecture as an explorable hierarchy and a left-to-right Sankey, making visible where value concentrates as it flows down the layers.

WHAT YOU'LL SEE

- Tree view: expand any node to drill down every level
- Value Flow view: Stakeholder → Value Proposition → Value Stream → Value Stage → Capability → AI Use-Case → Data Product
- A 'Weight by' selector (AI value score or KPI roll-up) and a start-from stakeholder picker

HOW TO USE IT

- Click a node to focus it and read its explanation and value
- Switch to Value Flow to see where value concentrates
- If you've made a selection elsewhere, the flow narrows to just that path

Open Graph →

Value Streams

Value Streams — Ownership & Governance `value_streams.html`

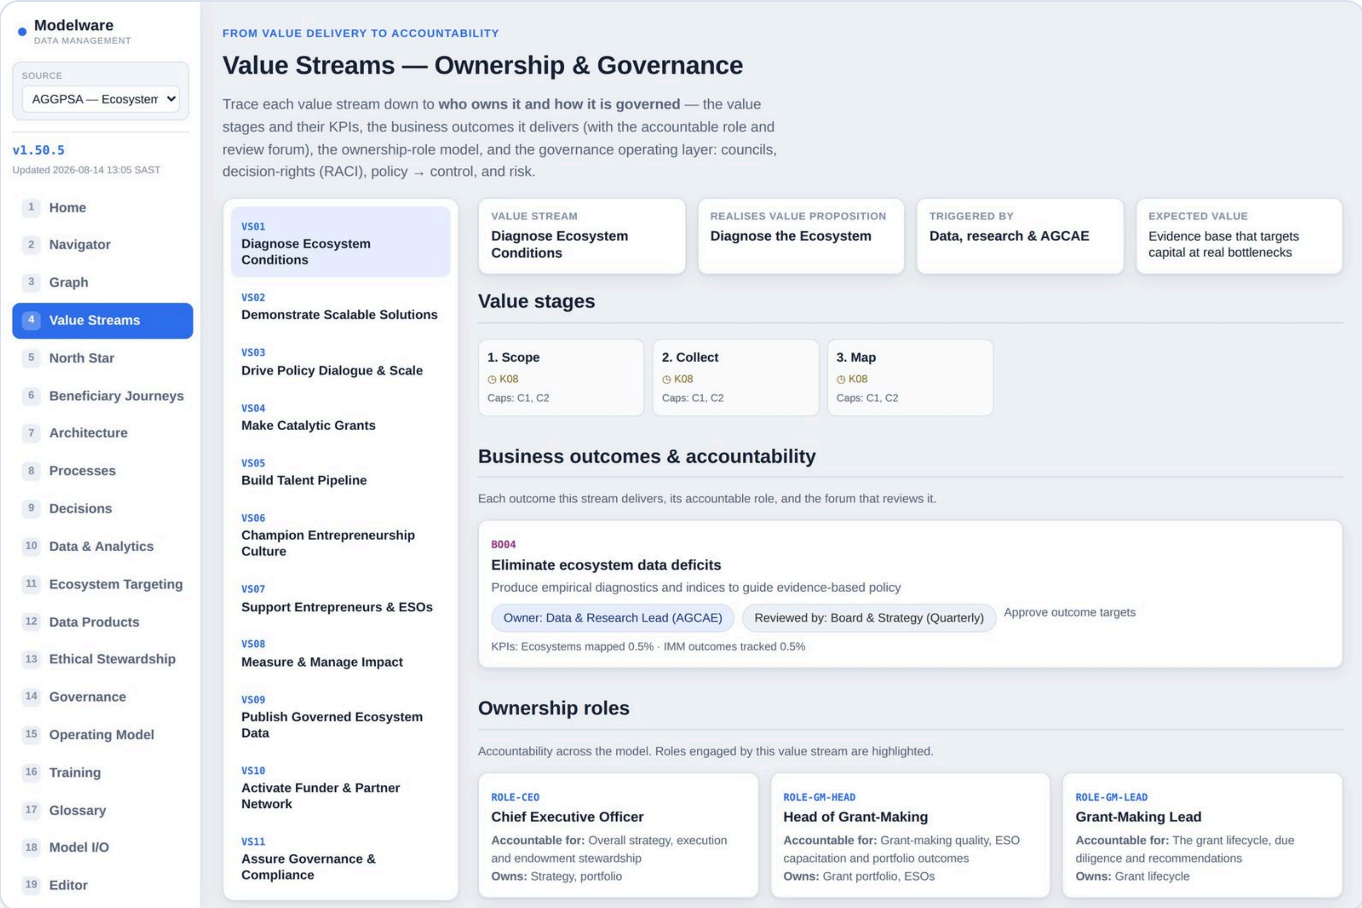

Trace each value stream down to who owns and governs it — value stages and KPIs, the business outcomes it delivers with their accountable roles and forums, the ownership-role model, and the governance operating layer.

> **ARCHITECTURE MESSAGE**
>
> The value layer with its governance overlay. Presents how the organisation creates value (streams → stages → KPIs → outcomes) and, laid over it, who owns and governs each part — accountable roles, councils, decision-rights RACI, policy → control and the risk register.

WHAT YOU'LL SEE

- A value-stream picker
- Value stages with primary KPIs
- Business outcomes with accountable roles and review forums
- Councils, decision-rights RACI, policy → control, and the risk register

HOW TO USE IT

- Pick a value stream on the left
- Read down from stages and KPIs to outcomes and owners
- Use the governance layer to confirm accountability

Open Value Streams →

North Star

North Star & Metric Tree `north_star.html`

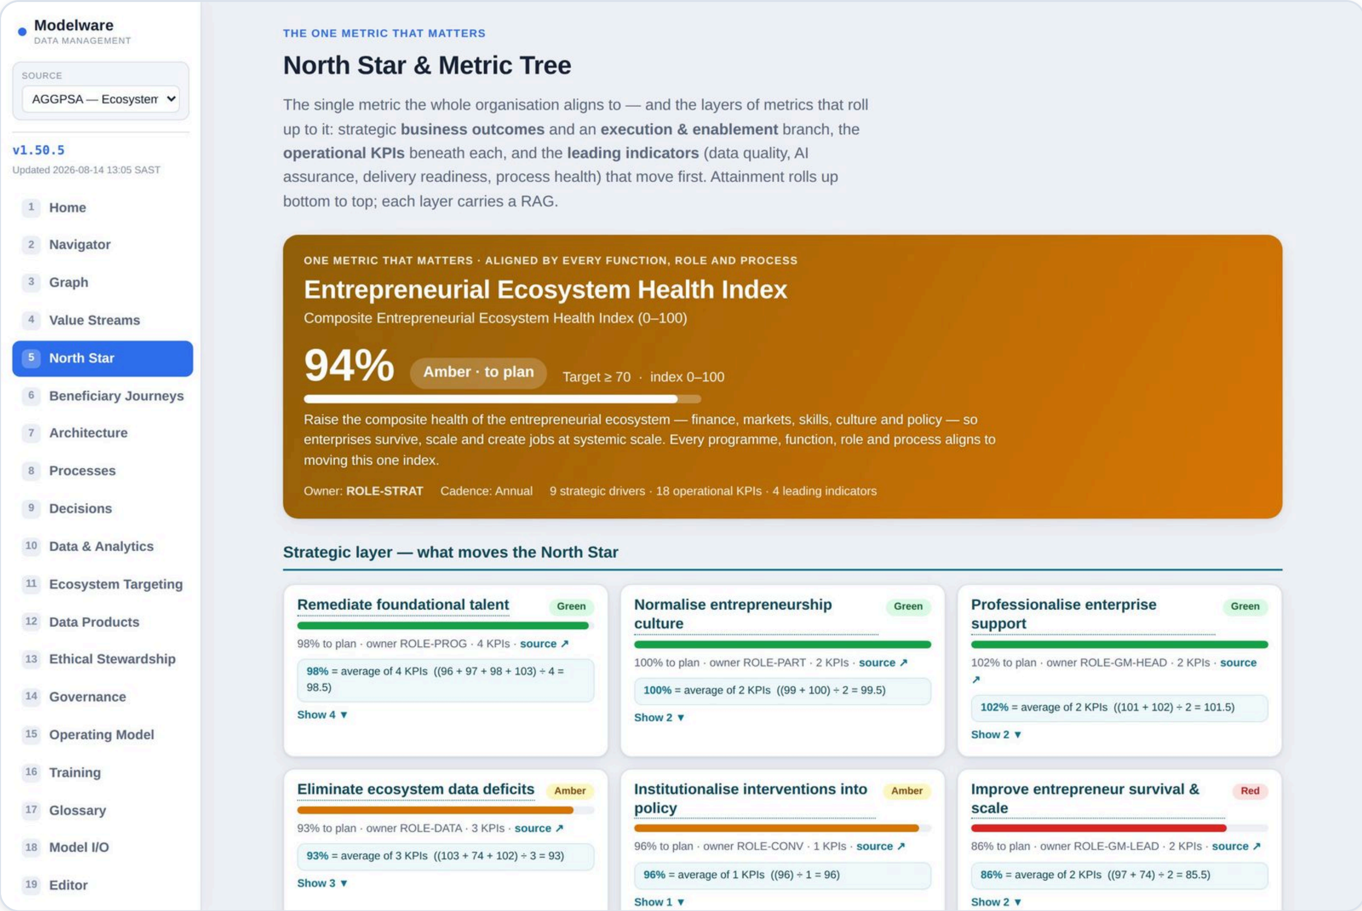

The One Metric That Matters (OMTM) and the layers of metrics that roll up to it — the single North Star the whole organisation aligns to, its strategic business-outcome drivers plus an execution & enablement branch, the operational KPIs beneath each, and the leading indicators (data quality, AI assurance, delivery readiness, process health) that move first. It also carries the AI model & agent eval dashboard.

> **ARCHITECTURE MESSAGE**
>
> The measurement layer. Presents the single North Star and the metric tree that rolls up to it — strategic drivers, operational KPIs and the leading enabler indicators (data quality, AI assurance, readiness) — so every layer of the architecture is tied to one number, alongside the AI model & agent assurance evals.

WHAT YOU'LL SEE

- The North Star with its attainment-to-plan and RAG
- Strategic drivers → operational KPIs (expandable), and the leading enabler indicators
- An AI eval dashboard — pass rate, and every model/agent eval with metric, threshold, result and Pass/Warn/Fail

HOW TO USE IT

- Read the North Star first — everything rolls up to it
- Open a strategic driver to see the KPIs beneath it
- Filter the evals to the failing/warning ones to see where AI assurance needs attention

Open North Star →

Beneficiary Journeys

Beneficiary Journeys & CX `customer_journey.html`

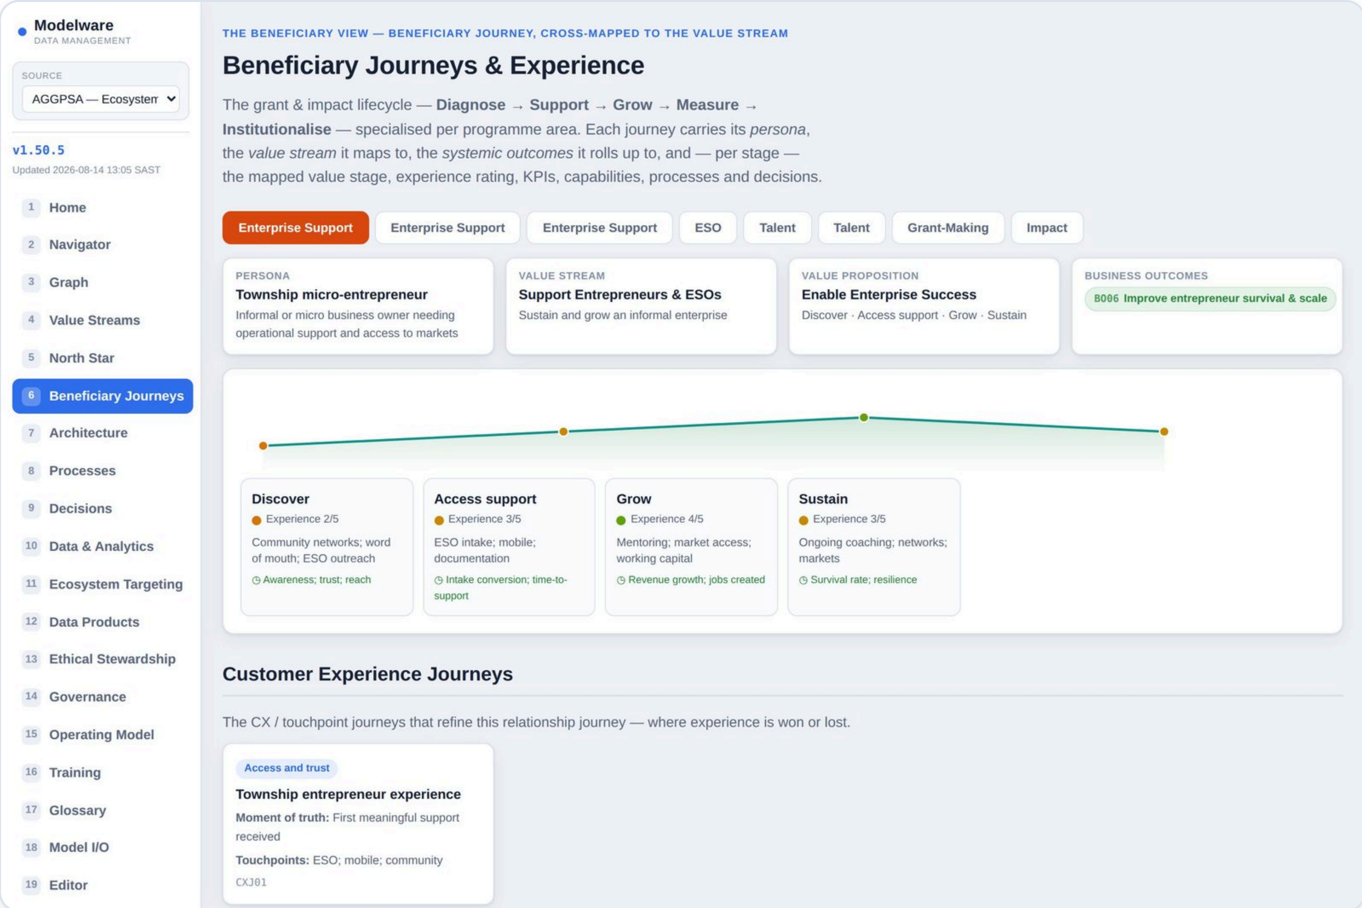

The beneficiary journey-map explorer across the grant & impact lifecycle, per programme area, with the catalytic-leverage governance view and the impact & learning operating model.

> **ARCHITECTURE MESSAGE**
>
> The experience layer, cross-mapped to the value layer. Presents the outside-in beneficiary/customer view of the very same value the streams create — keeping the human experience aligned to the internal architecture stage-by-stage.

WHAT YOU'LL SEE

- Journey tabs per programme area (Enterprise Support, Talent, Grant-Making, Impact)
- Each stage with its experience score, touchpoints, capabilities, processes and KPIs
- The mapped value stream and systemic outcomes

HOW TO USE IT

- Pick a programme-area tab
- Read left-to-right along the journey stages
- Click a KPI or stage to see detail

Open Beneficiary Journeys →

Architecture

Business Architecture `business_architecture.html`

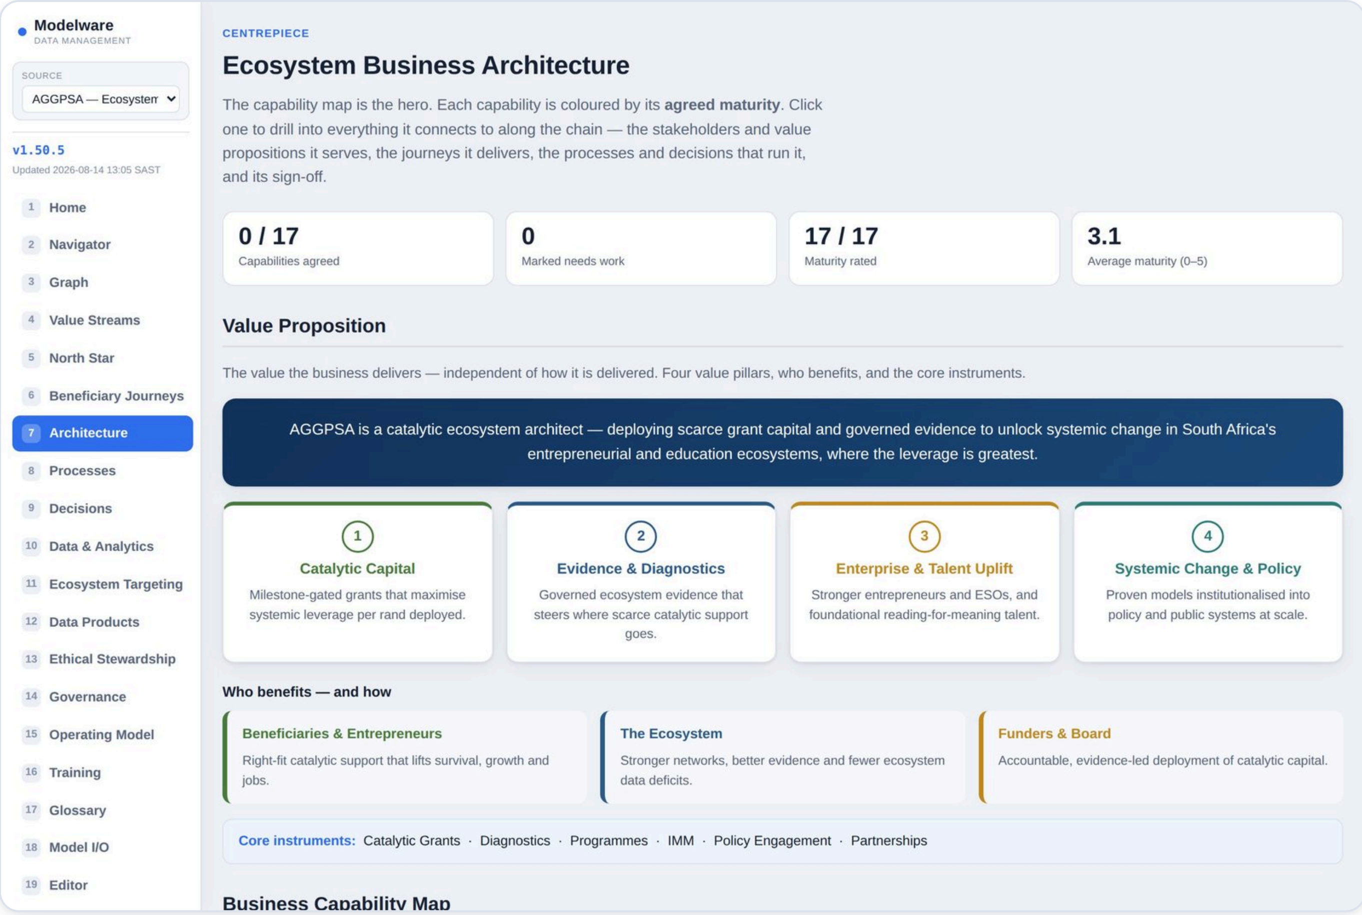

The capability centrepiece. Shows the capability map under a chosen value proposition, with value-stream heat and an Attention/Maturity toggle, plus the full capability grid for sign-off.

> **ARCHITECTURE MESSAGE**
>
> The capability layer — the centrepiece. Presents what the organisation must be able to do (capabilities), framed by the value proposition and coloured by maturity/attention, as the reusable building blocks every stream, process and AI investment draws on.

WHAT YOU'LL SEE

- A capability map coloured by heat
- An Attention vs Maturity toggle
- A drill-in panel for any capability

HOW TO USE IT

- Choose a value proposition to frame the map
- Toggle between Attention and Maturity
- Click a capability to drill in and sign it off

Open Architecture →

Processes

Business Processes (SIPOC) `business_process.html`

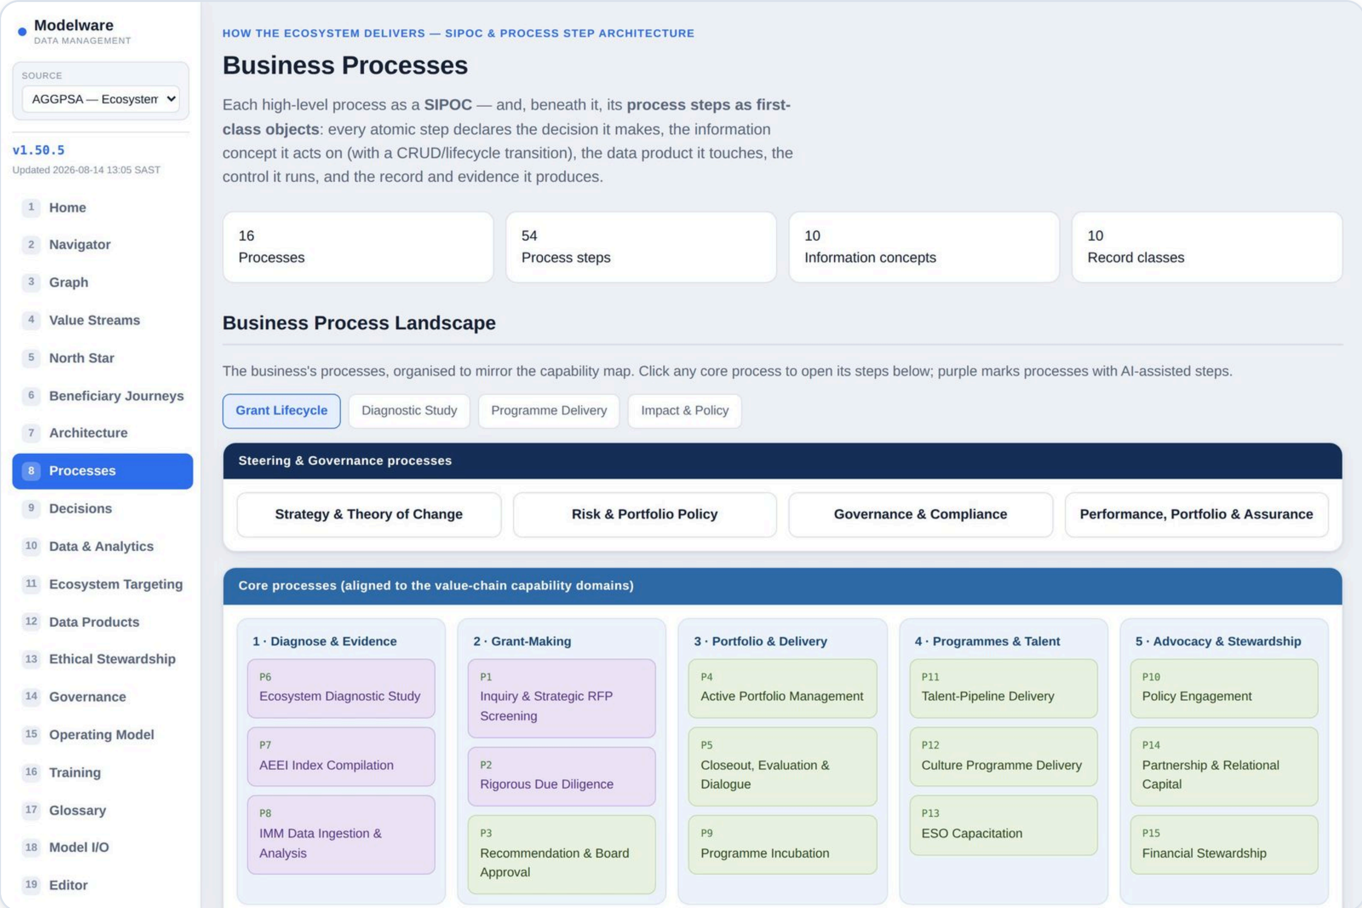

Every high-level process as a SIPOC — Suppliers, Inputs, Process, Outputs, Customers — with its capabilities and participants, traced back to stakeholders.

> **ARCHITECTURE MESSAGE**
>
> The process layer, as SIPOC. Presents how each capability is actually performed — Suppliers, Inputs, Process, Outputs, Customers — linking the abstract capability down to real, ownable work with clear hand-offs.

WHAT YOU'LL SEE

- A SIPOC card per process
- The capabilities each process realises
- Participants and the stakeholders served

HOW TO USE IT

- Scan the SIPOC to understand a process at a glance
- Follow the capability links up to the architecture
- Use it to confirm ownership and hand-offs

Open Processes →

Decisions

Decision Models `decisions.html`

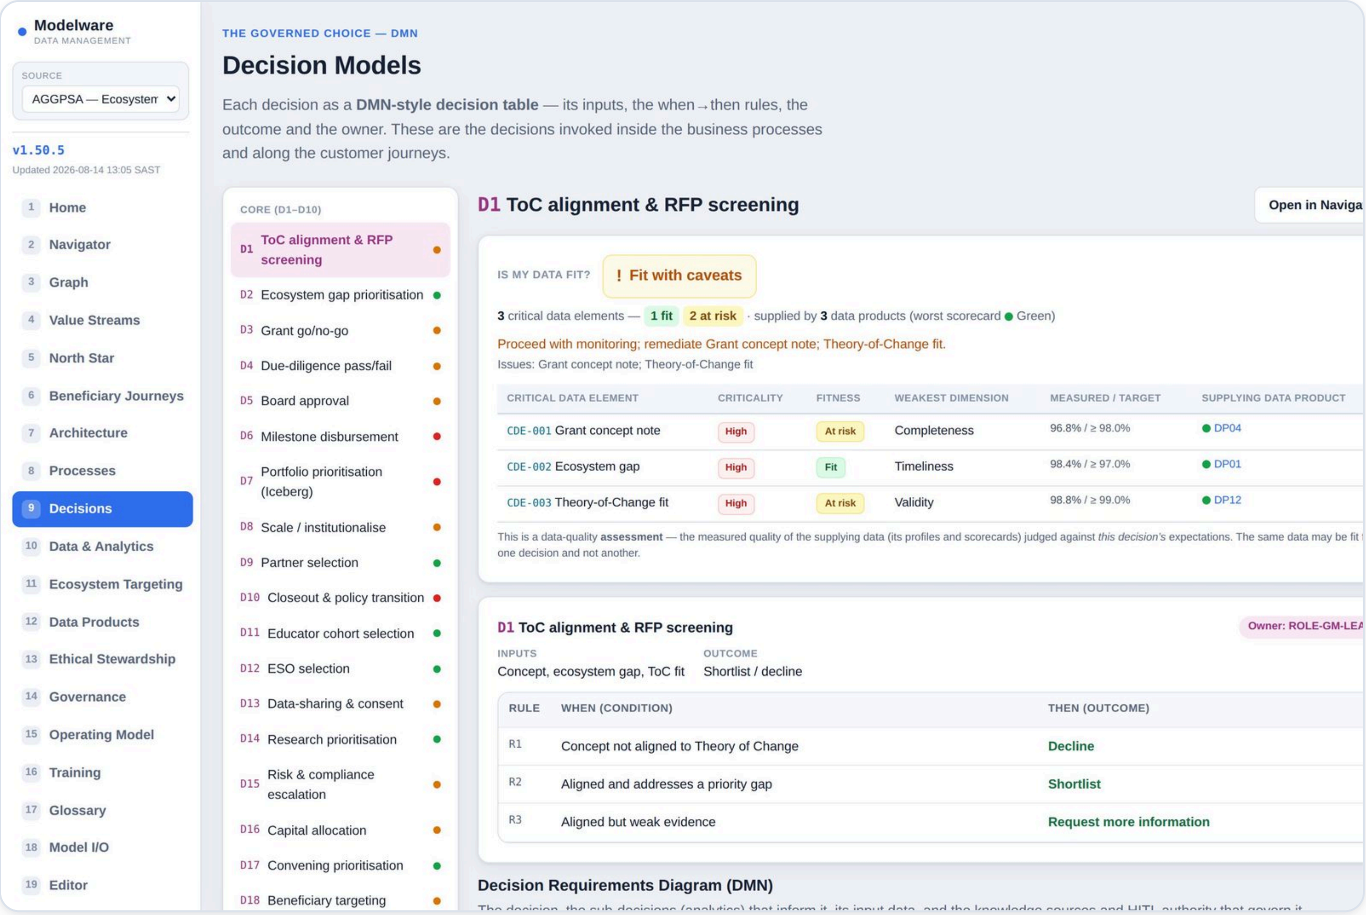

Every decision as a DMN-style decision table — inputs, when→then rules, the outcome and its owner — with where each one is used.

ARCHITECTURE MESSAGE

The decision layer. Presents the specific judgements made inside processes as governed decision tables — the exact points where AI attaches to the business and where accountability for an outcome is exercised.

WHAT YOU'LL SEE

- A decision table per decision
- Input conditions and the resulting outcome
- The owner and the processes/use-cases that call it

HOW TO USE IT

- Read the rules top-to-bottom
- Check the owner for accountability
- Trace a decision to the AI use-case that automates it

Open Decisions →

Data & Analytics

Ecosystem Analytics `ai_usecases.html`

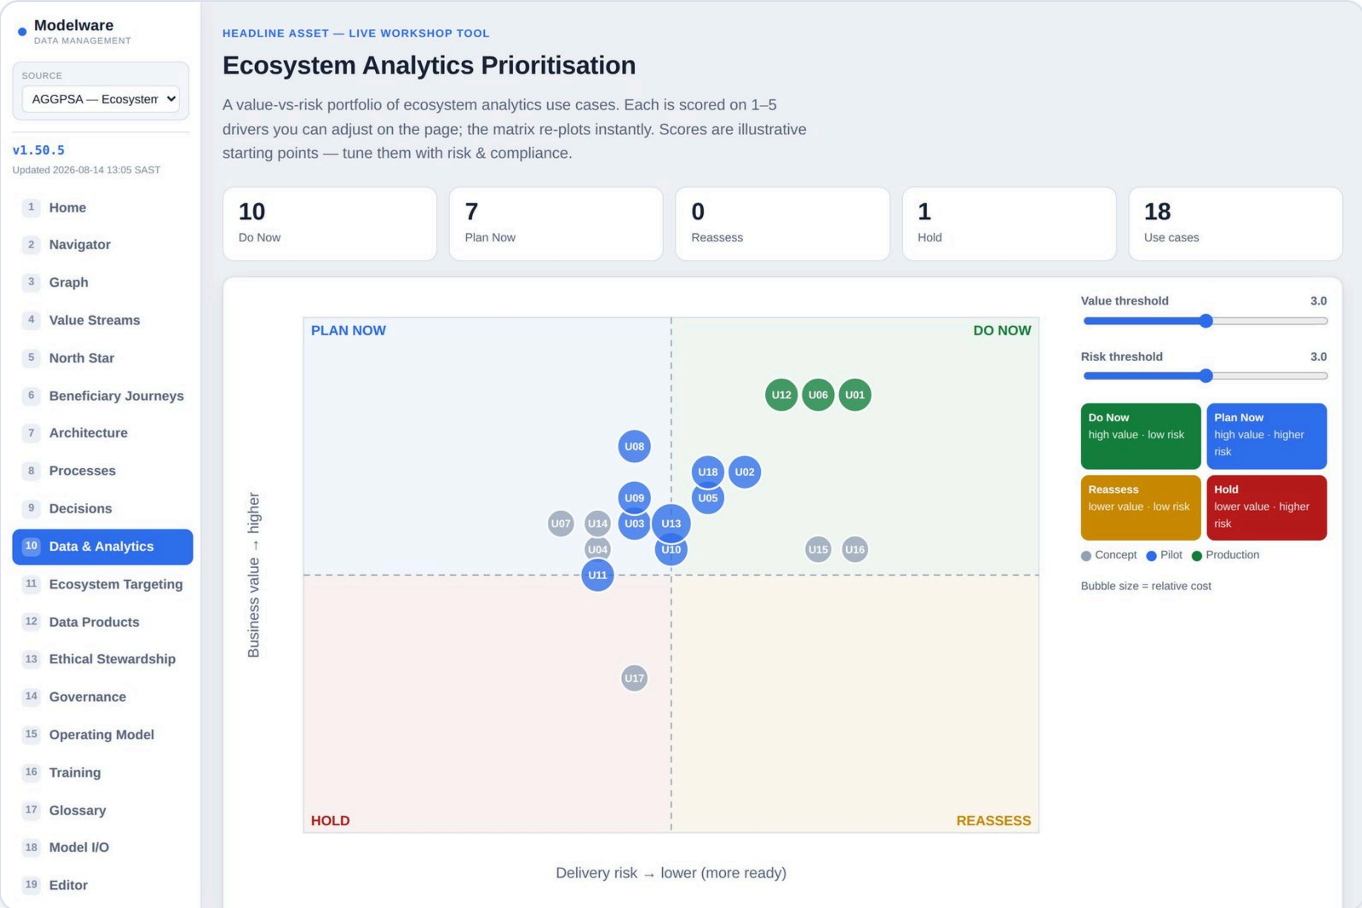

A value-vs-risk portfolio of ecosystem analytics use cases, with live 1–5 driver scoring and an Excel round-trip.

> **ARCHITECTURE MESSAGE**
>
> The AI-initiative layer, as a governed portfolio. Presents every AI use-case scored on value versus risk and assessed for readiness across six dimensions, so AI is prioritised as an evidenced portfolio anchored to decisions, outcomes and data — not sold on hype.

WHAT YOU'LL SEE

- A value-vs-risk bubble chart of every use-case
- Live driver sliders (systemic, evidence, leverage, reach, efficiency; data, model, equity, safeguarding, adoption)
- Export/import of the scoring model

HOW TO USE IT

- Read the portfolio: top-right = high value, watch the risk axis
- Adjust drivers to test scenarios
- Export to Excel, score offline, import back

Open Data & Analytics →

Ecosystem Targeting

Ecosystem Targeting `hyperpersonalisation_cdp.html`

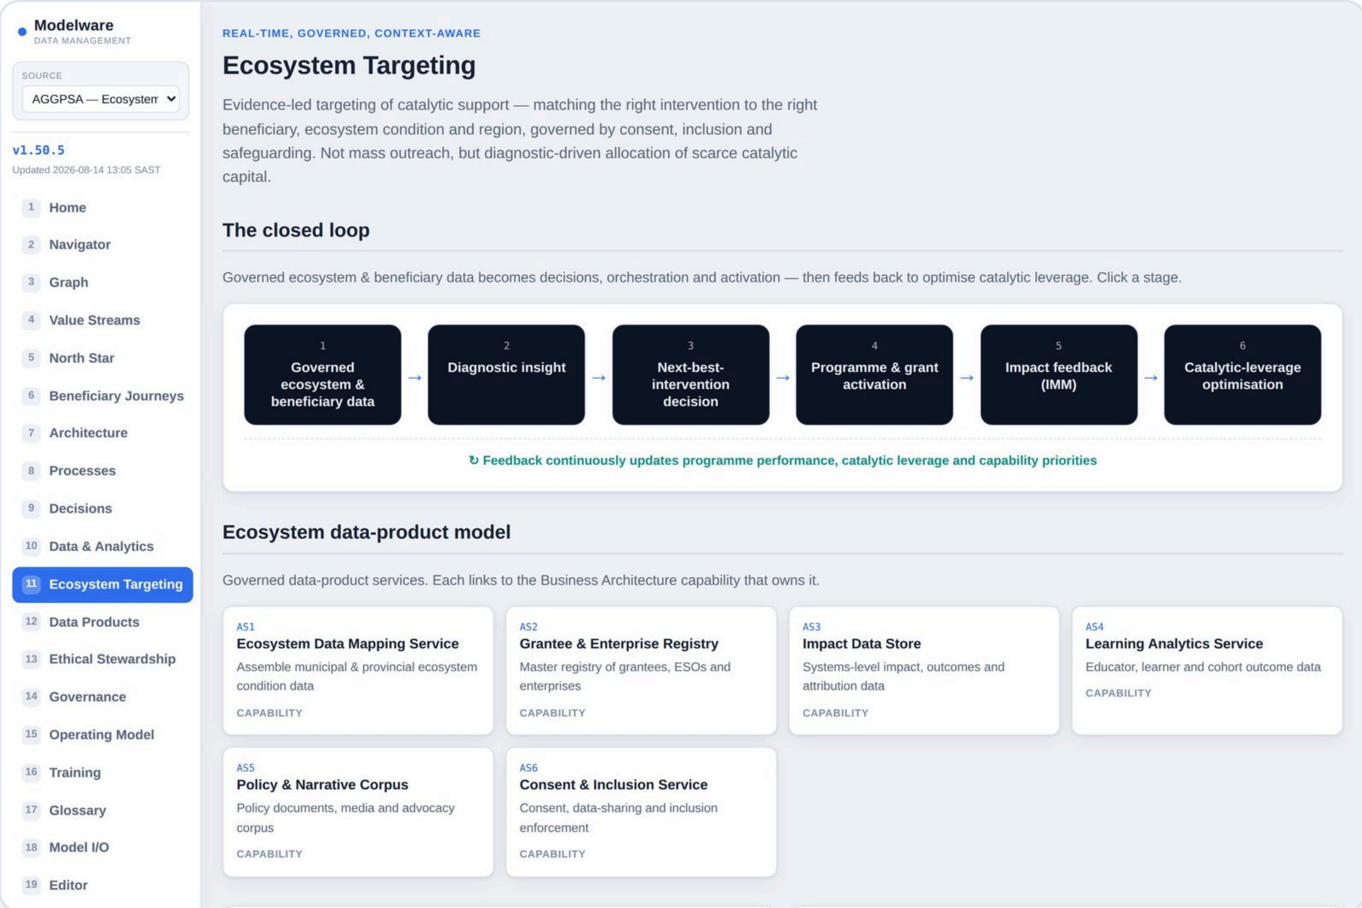

The closed-loop ecosystem targeting engine and the data-product model — the data, models and governance that make evidence-led targeting work.

> **ARCHITECTURE MESSAGE**
>
> The activation layer. Presents the closed-loop targeting/personalisation engine and its data-product model — how governed data and models are turned into the next best action, with the governance wrapped around it.

WHAT YOU'LL SEE

- The targeting closed loop
- The ecosystem data-product service model
- The governance guardrails around it

HOW TO USE IT

- Follow the loop from signal to action to measurement
- See which services each use-case needs
- Use it to frame the data and model dependencies

Open Ecosystem Targeting →

Data Products

Data Products `data_products.html`

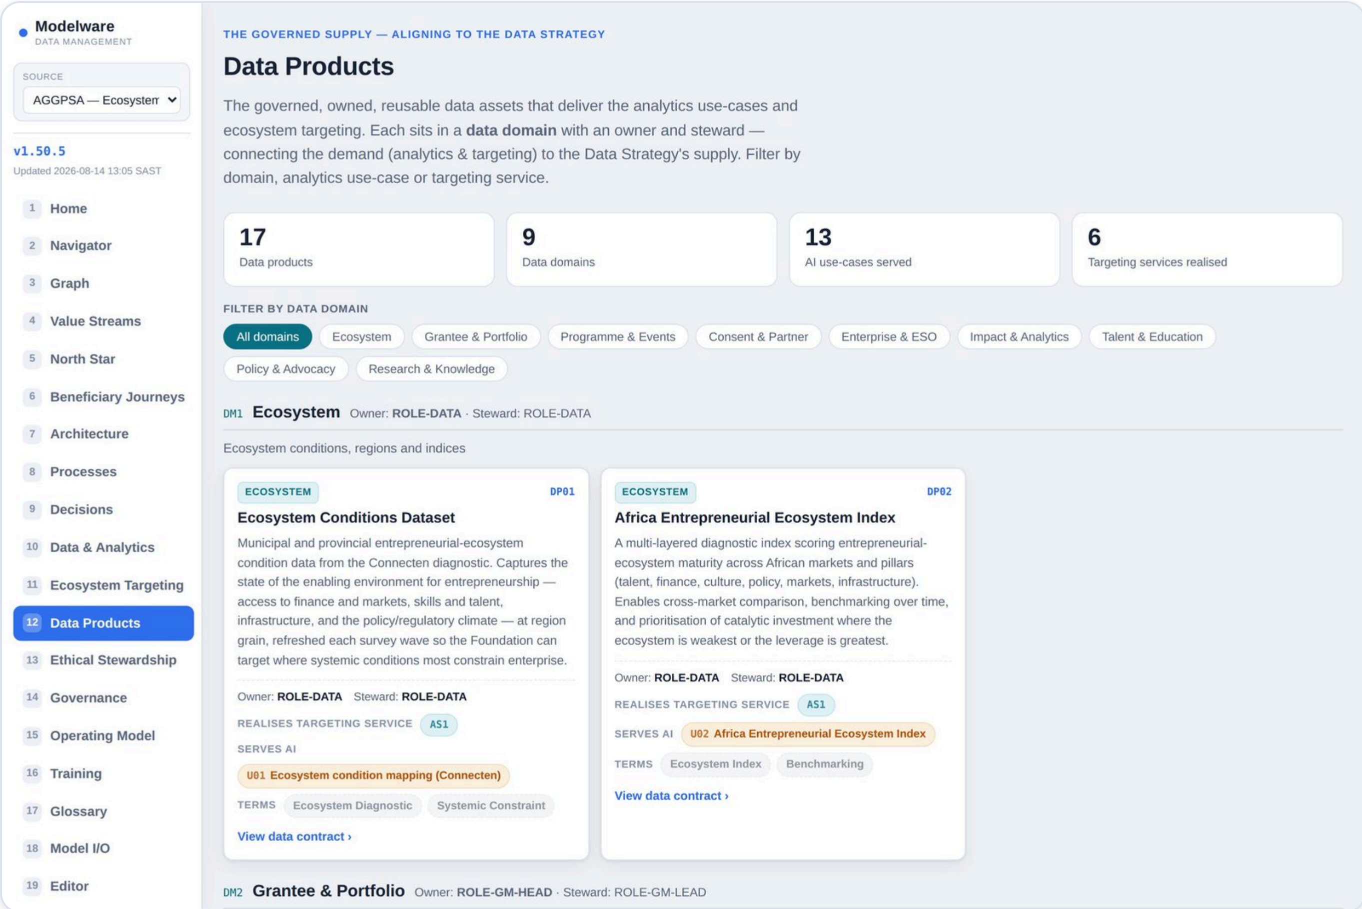

The governed, owned data products that deliver the analytics use-cases and ecosystem targeting — organised by data domain, each with an owner, steward and data contract.

> **ARCHITECTURE MESSAGE**
>
> The data-product layer. Presents the governed, owned data assets the AI and analytics actually run on — organised by domain, each with an owner, steward and data contract — making the data dependency of every initiative explicit and accountable.

WHAT YOU'LL SEE

- Data products grouped by domain
- Owner and steward per product
- A data-contract panel (schema, classification, sources)

HOW TO USE IT

- Filter by domain
- Click a product to open its data contract
- Trace a product up to the use-cases it serves

Open Data Products →

Ethical Stewardship

Ethical Stewardship — the orientation layer `ethical_stewardship.html`

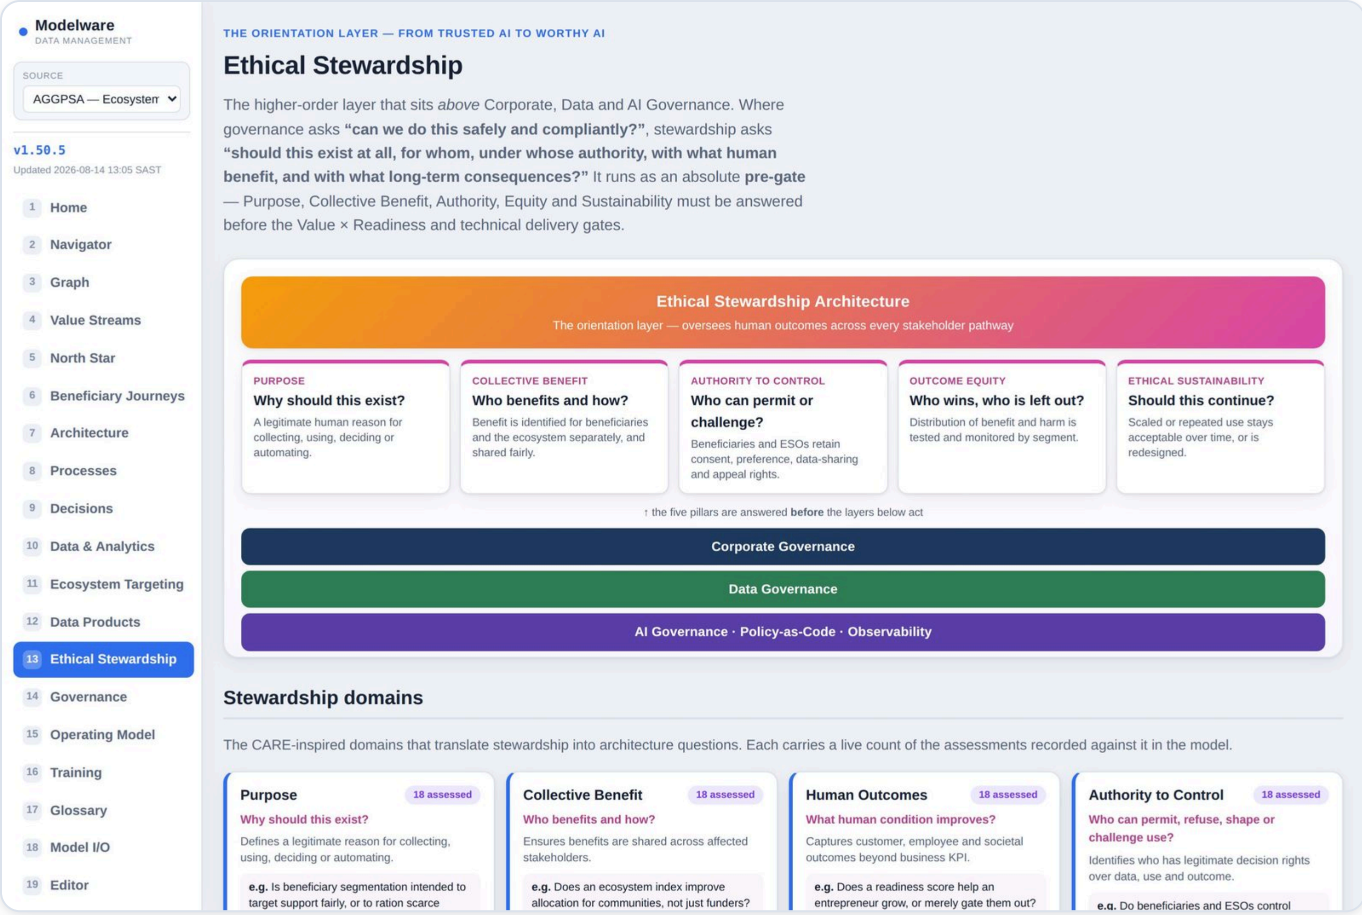

The ethical-stewardship layer — human outcomes, equity and inclusion, collective benefit and the stewardship roles that hold them — so the model shows not just what is efficient but what is responsible and fair.

> **ARCHITECTURE MESSAGE**
>
> The orientation & values layer. Presents the ethical-stewardship principles that sit above the whole model, so the architecture shows not only what is efficient but what is responsible, fair and human-centred.

WHAT YOU'LL SEE

- Human-outcome and equity/inclusion elements
- Collective-benefit statements
- The stewardship roles accountable for them

HOW TO USE IT

- Read the human outcomes each initiative is meant to protect
- Check equity and inclusion considerations
- Confirm who stewards responsible outcomes

Open Ethical Stewardship →

Governance

Governance & Responsible AI `governance_responsible_ai.html`

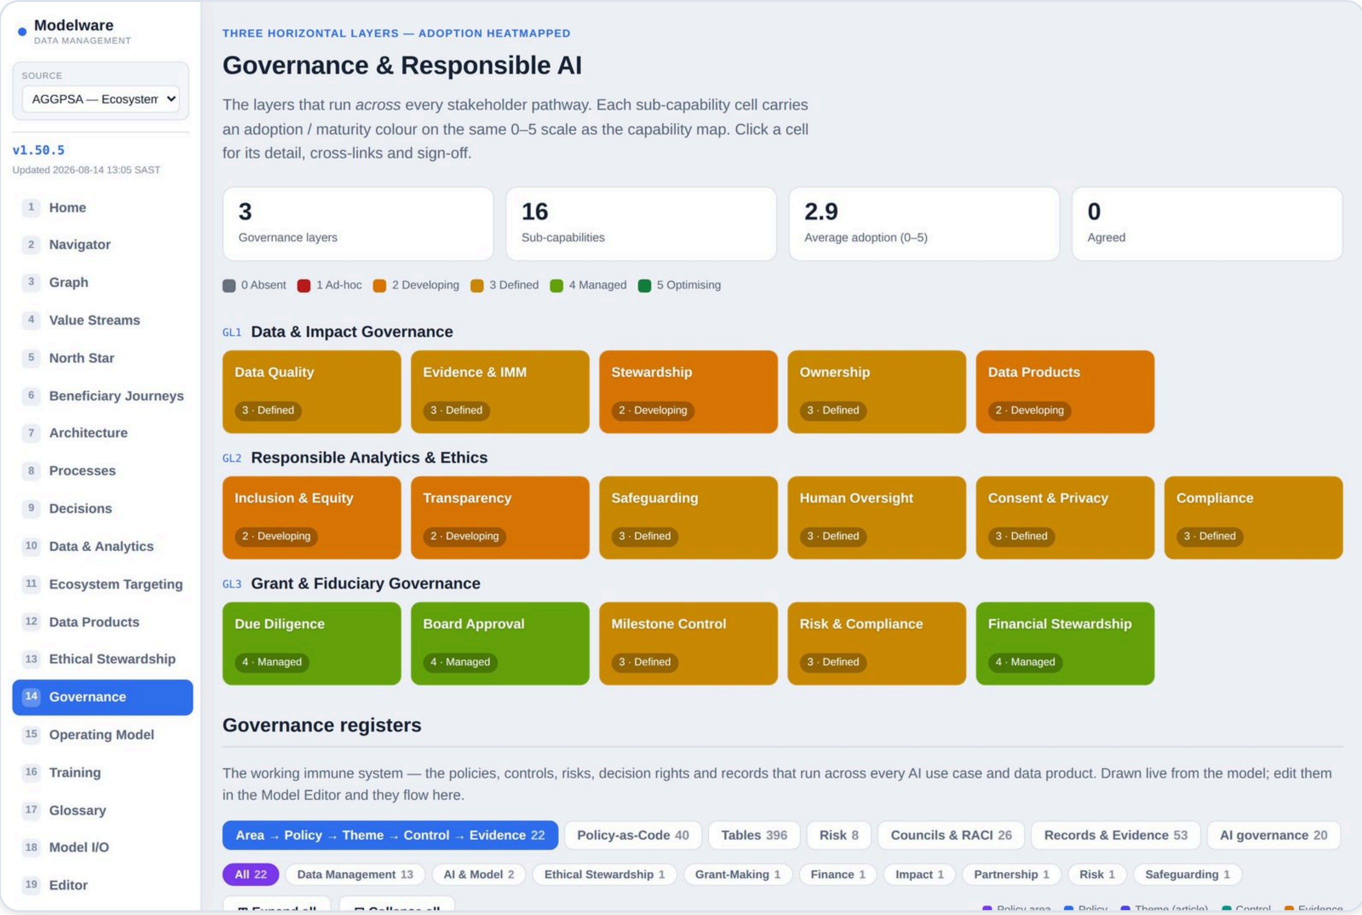

A layered heatmap of Data Governance, Responsible AI and Change & Adoption, with a success-measure scorecard.

ARCHITECTURE MESSAGE

The governance layer, as measurable maturity. Presents Data Governance, Responsible AI and Change & Adoption as a heatmap, carrying the Area → Policy → Theme → Control → Evidence chain and a success-measure scorecard.

WHAT YOU'LL SEE

- A layered governance heatmap
- Responsible-AI controls
- A success-measure scorecard

HOW TO USE IT

- Read the heat to spot weak layers
- Check the scorecard against targets
- Use it to brief governance forums

Open Governance →

Operating Model

Operating Model `operating_model.html`

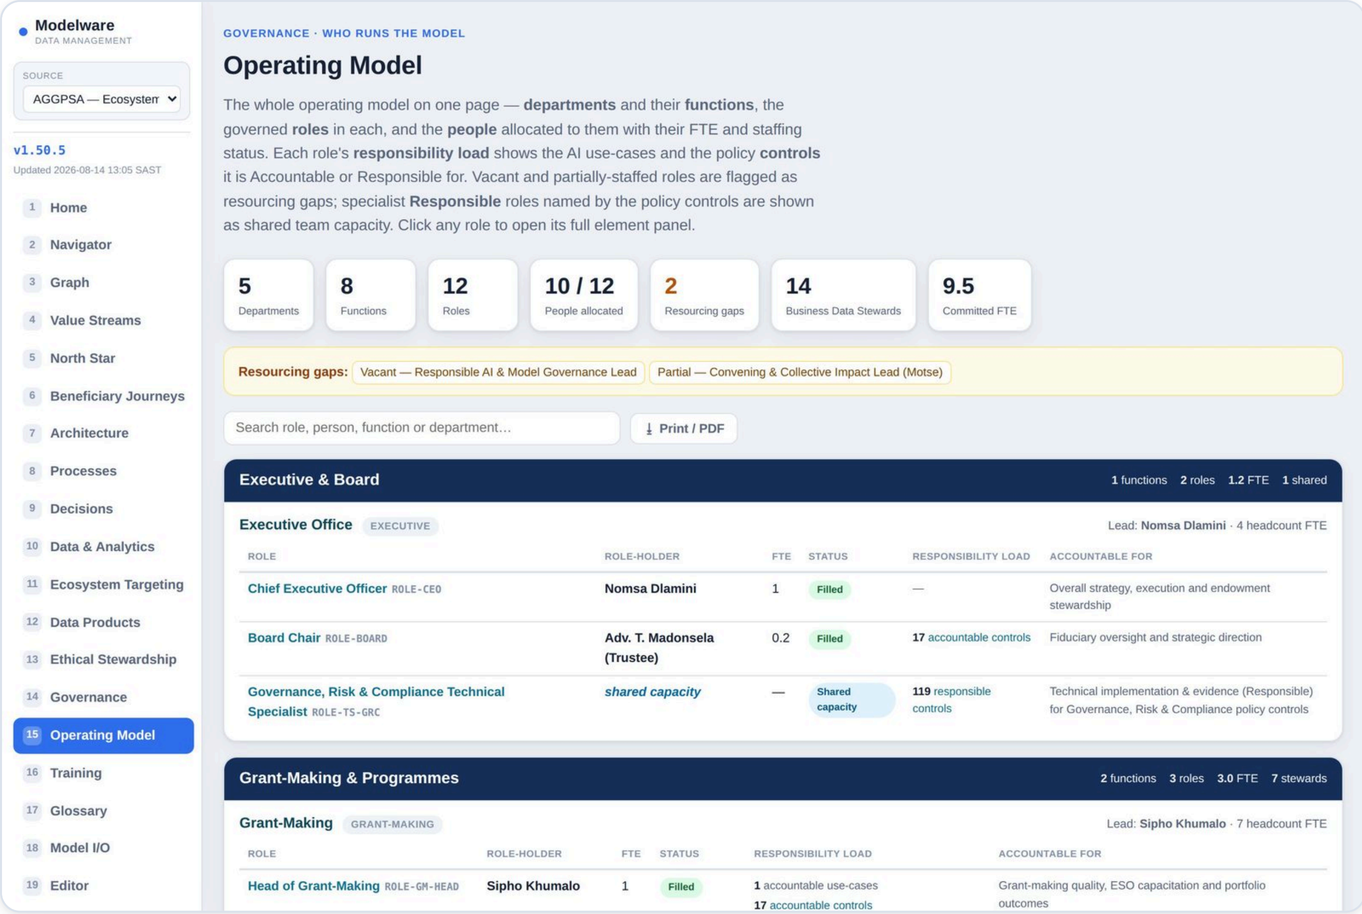

The whole operating model on one page — departments and their functions, the governed roles in each, and the people allocated to them with FTE and staffing status. This is what grounds People readiness: a vacant accountable role is a real resourcing gap.

> **ARCHITECTURE MESSAGE**
>
> The people & organisation layer. Presents the departments, functions, governed roles and the people in them — grounding the model in who actually runs it, each role's use-case and policy-control responsibility load, the Business Data Steward network, and where the resourcing gaps are.

WHAT YOU'LL SEE

- Departments → functions → roles → allocated people, with FTE and Filled/Partial/Vacant status
- A summary of roles, people allocated, resourcing gaps and committed FTE
- Each role's use-case load (how many it is accountable/responsible for)

HOW TO USE IT

- Scan a department to see how it is staffed
- Spot resourcing gaps (vacant / partially-staffed roles) at a glance
- Click a role to open its full element panel

Open Operating Model →

Training

Training & Adoption `training_adoption.html`

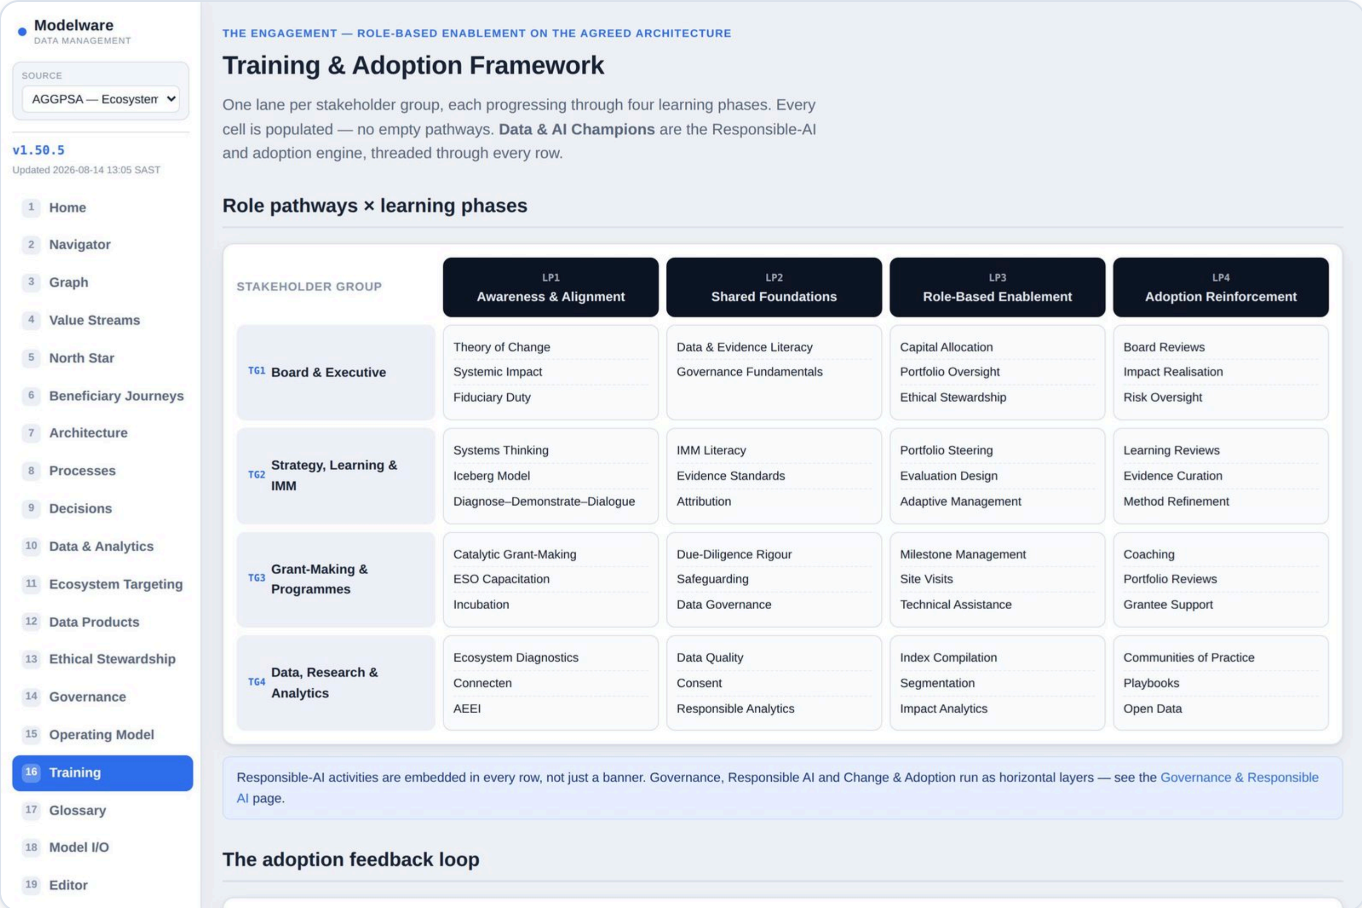

Role pathways across four learning phases, with Champions as the Responsible-AI engine and the adoption feedback loop.

> **ARCHITECTURE MESSAGE**
>
> The adoption & change layer. Presents how capability is built into people — role learning pathways and the Champions network — so the architecture is operationalised and sustained, not merely documented.

WHAT YOU'LL SEE

- Learning pathways per role
- Four phases from awareness to mastery
- The Champions model and feedback loop

HOW TO USE IT

- Find your role's pathway
- Follow the four phases in order
- See how Champions carry Responsible-AI into the business

Open Training →

Glossary

Glossary settings `glossary_settings.html`

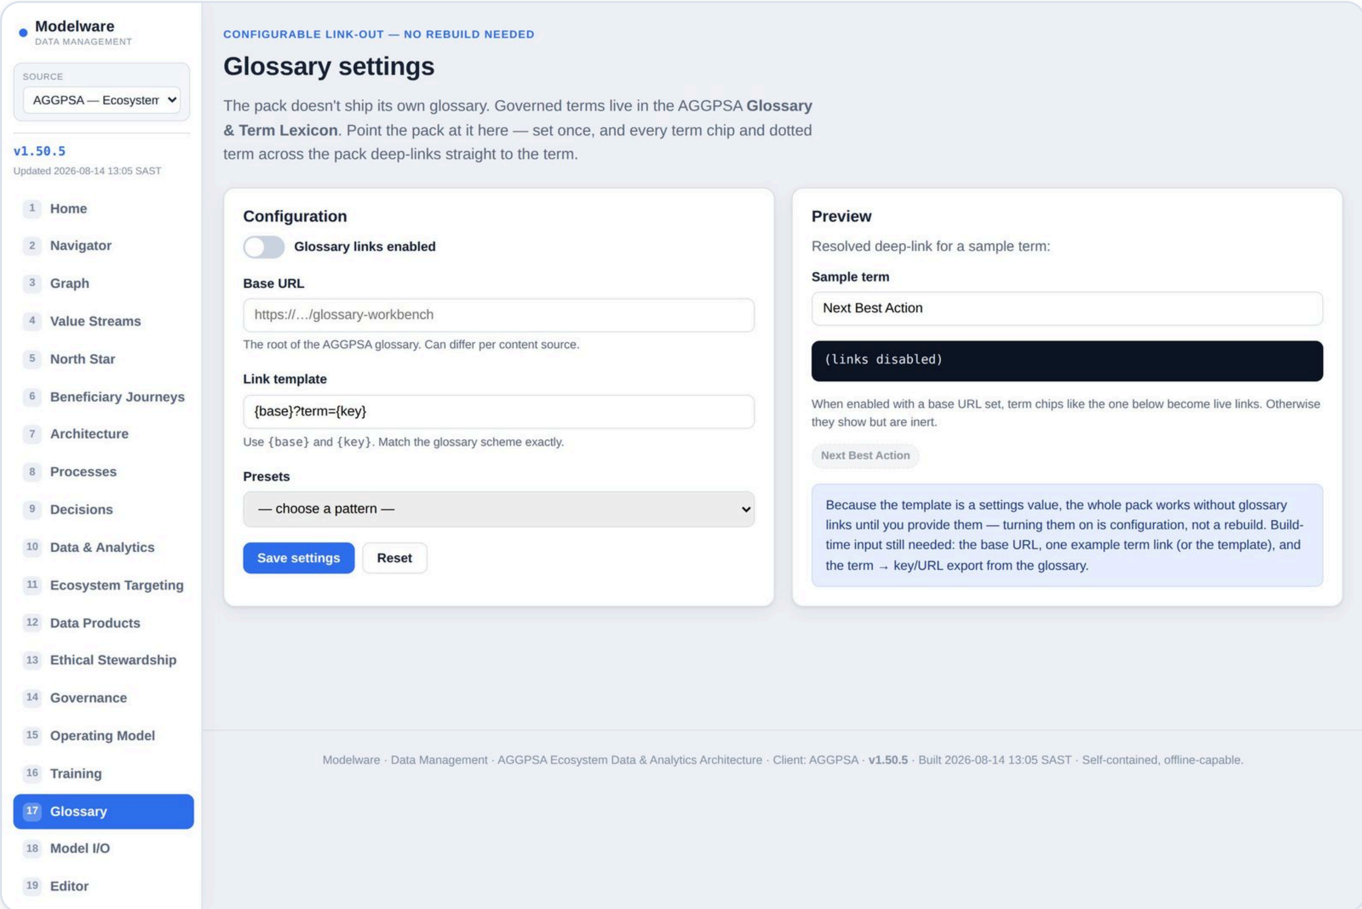

Point the pack at the AGGPSA glossary — base URL and link template — so governed terms deep-link out to their definitions.

ARCHITECTURE MESSAGE

The semantic layer's gateway. Connects the pack's governed terms to the external business glossary so every element resolves to one agreed definition — the shared vocabulary that lets the whole model's numbers reconcile.

WHAT YOU'LL SEE

- A base-URL setting
- A link-template setting
- A live preview of a resolved term link

HOW TO USE IT

- Set the glossary base URL and template
- Save the settings
- Governed terms across the pack then deep-link to the glossary

Open Glossary →

Model I/O

Model Export / Import `model_export_import.html`

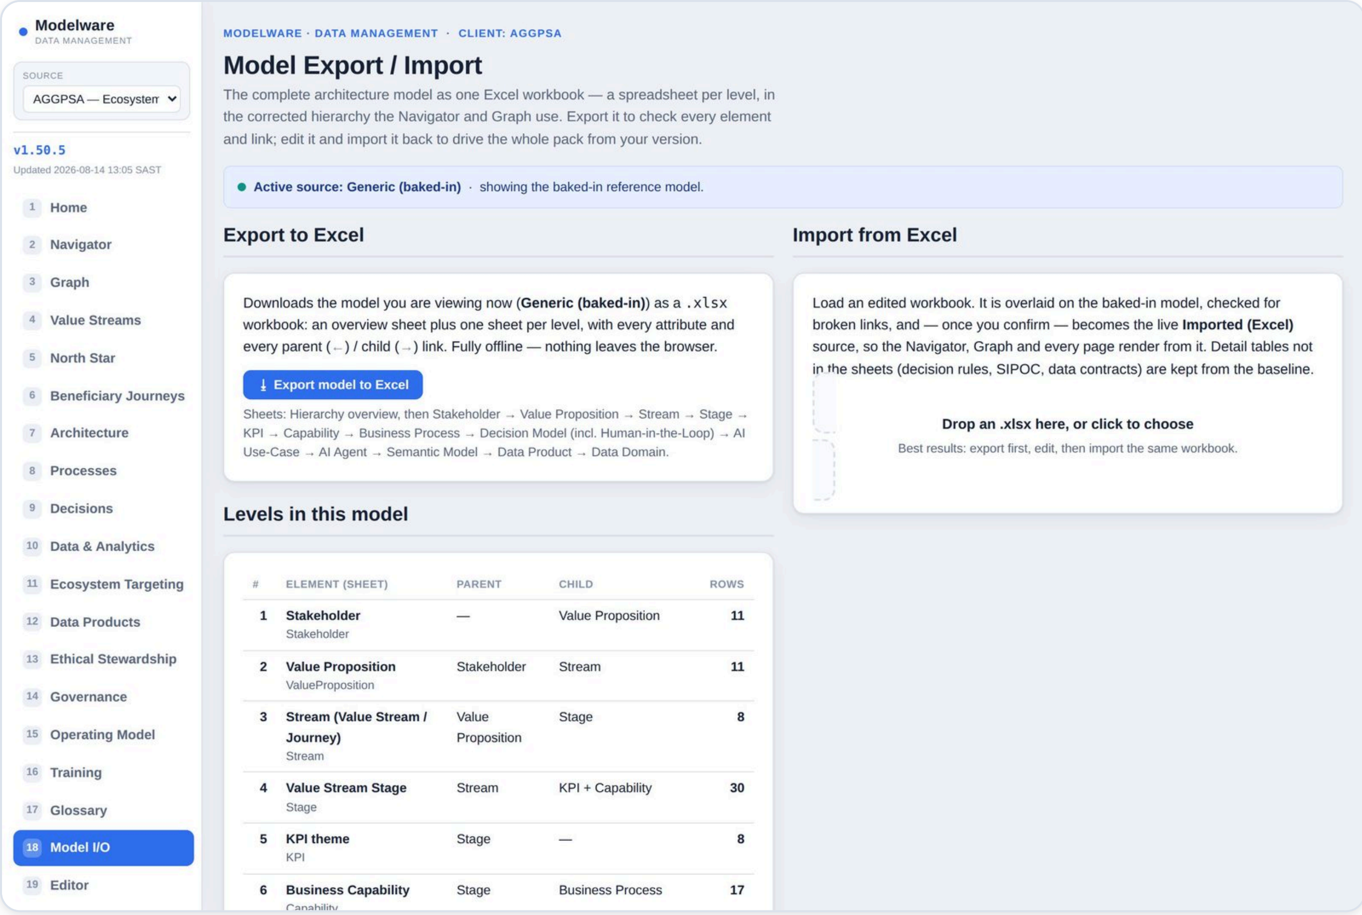

Export the complete model to Excel — one sheet per level, every attribute and link, with a referential-integrity check — and import an edited workbook back as a live source that drives every page.

ARCHITECTURE MESSAGE

The model-as-data control point. Presents the entire architecture as a round-trippable Excel workbook with a referential-integrity check — proving the model is a governed, verifiable dataset, not a picture.

WHAT YOU'LL SEE

- An Export to Excel action
- An Import from Excel action
- A referential-integrity validation summary

HOW TO USE IT

- Export to verify the model in a spreadsheet
- Edit and re-import to load it as the live 'Imported (Excel)' source
- Use the integrity check to catch broken links

Open Model I/O →

Editor

Model Editor `model_editor.html`

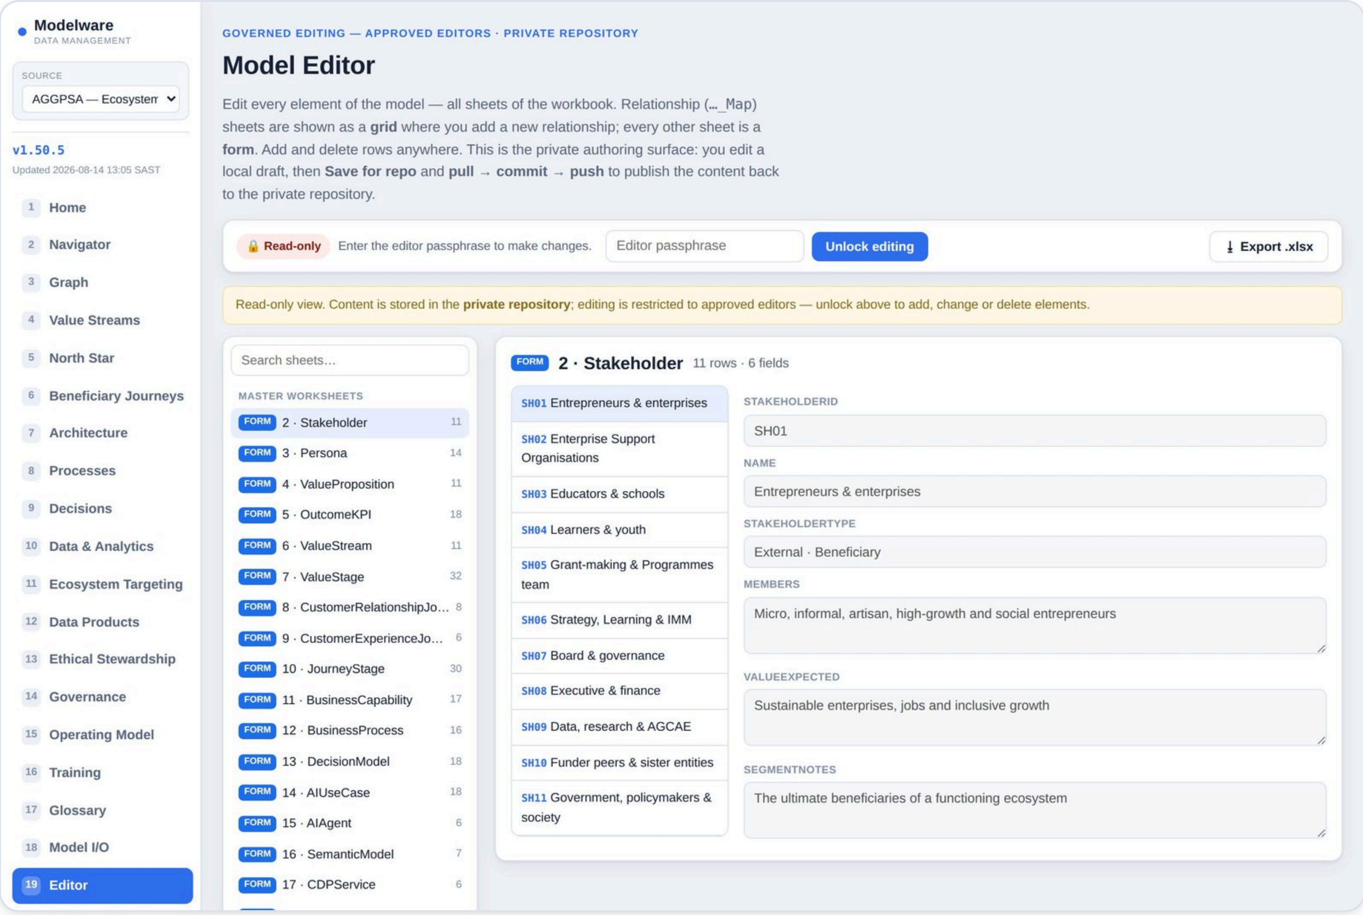

Edit every element of the model — all sheets. Relationship (…_Map) sheets are editable grids where you add a relationship; every other sheet is a form. Restricted to approved editors, and changes are published through the private repository.

ARCHITECTURE MESSAGE

The model-authoring layer. Presents every sheet of the architecture as editable forms and relationship grids under approval control — the governed way the model is changed and published to everyone.

WHAT YOU'LL SEE

- A passphrase gate (approved editors only)
- Grids for relationship sheets; forms for everything else
- Add / delete rows anywhere; Save for repo, Export .xlsx, Load published

HOW TO USE IT

- Unlock with the editor passphrase
- Edit a local draft in the browser
- Save for repo, then pull → commit → push to publish to everyone

Open Editor →

Manual

User Manual `user_manual.html`

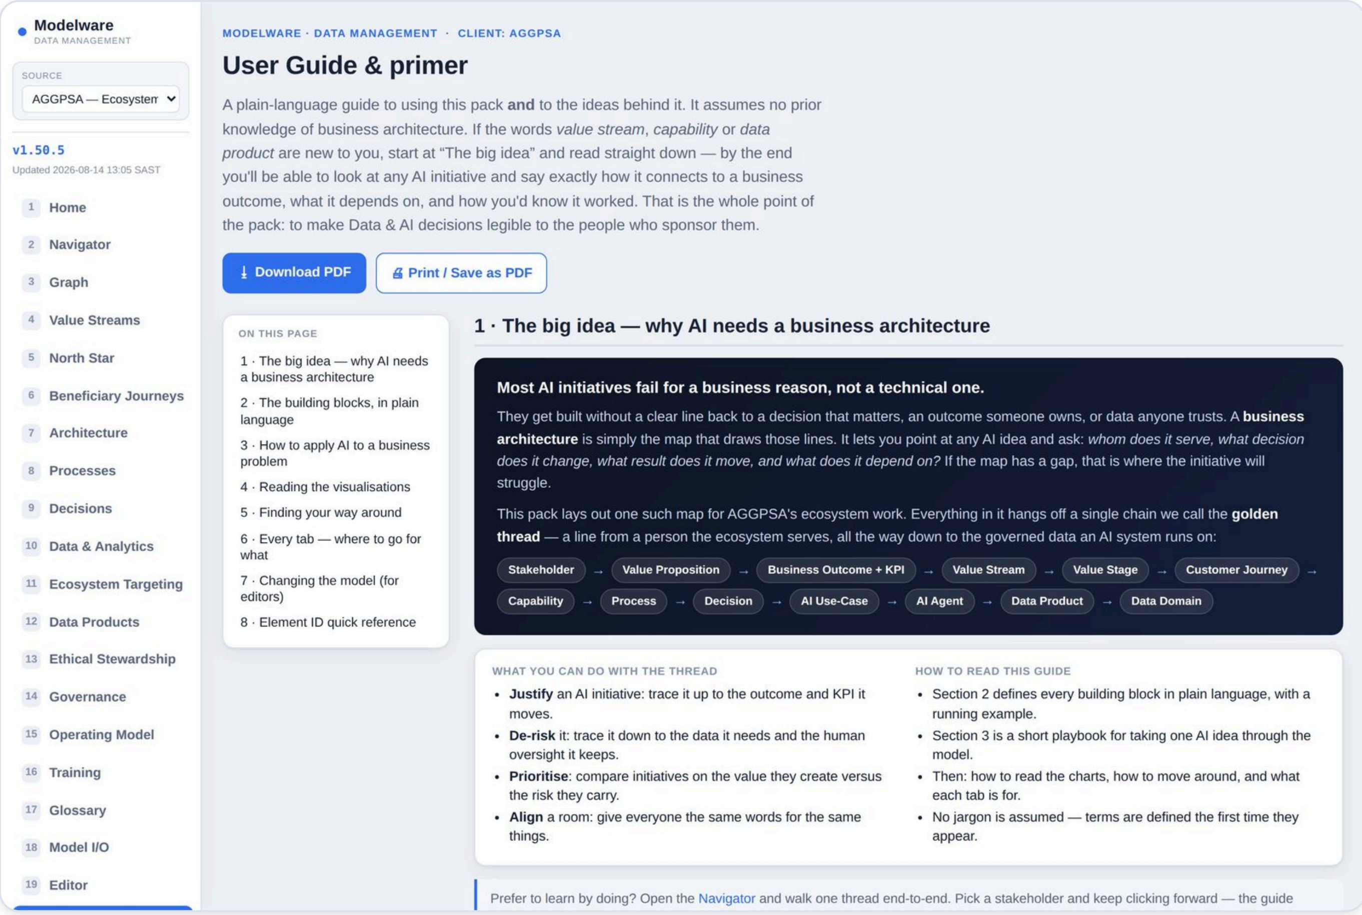

This manual. A living reference to every tab and the Navigator flow, generated from the same configuration the app runs on, so it stays in step with the pack.

> **ARCHITECTURE MESSAGE**
>
> The guide to the whole architecture. Explains the golden thread, defines every building block in plain language, and documents what each page is for and the architecture message it carries.

WHAT YOU'LL SEE

- A section for every tab
- The 13-step Navigator flow explained
- Download / print controls at the top

HOW TO USE IT

- Use the contents list on the left to jump to a topic
- Download the PDF to share offline or attach to onboarding

Critical Data

Critical Data Elements `critical_data_elements.html`

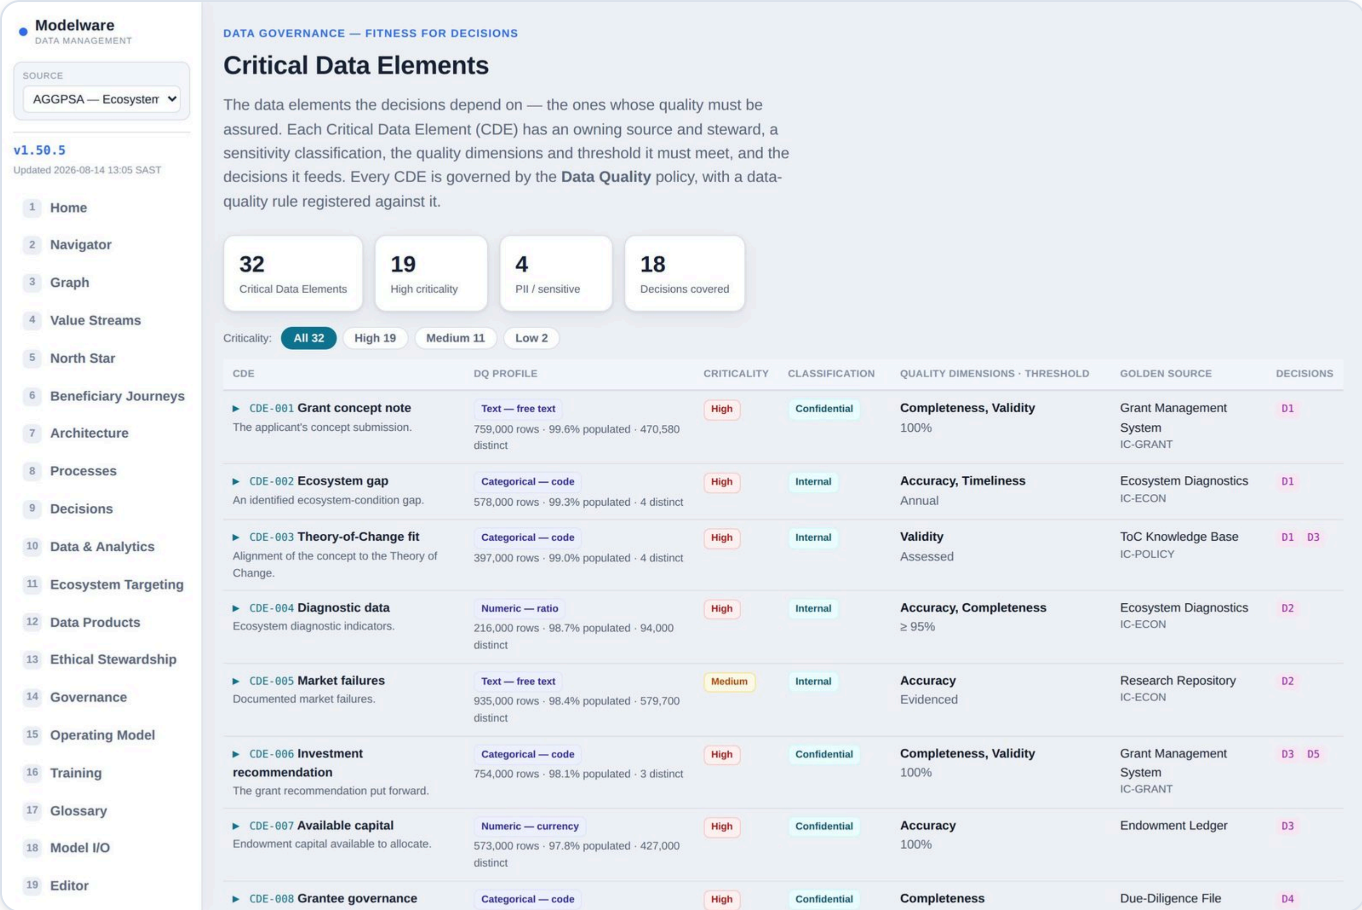

CDE	DQ PROFILE	CRITICALITY	CLASSIFICATION	QUALITY DIMENSIONS · THRESHOLD	GOLDEN SOURCE	DECISIONS
▶ CDE-001 **Grant concept note** The applicant's concept submission.	Text — free text 759,000 rows · 99.6% populated · 470,580 distinct	High	Confidential	**Completeness, Validity** 100%	Grant Management System IC-GRANT	D1
▶ CDE-002 **Ecosystem gap** An identified ecosystem-condition gap.	Categorical — code 578,000 rows · 99.3% populated · 4 distinct	High	Internal	**Accuracy, Timeliness** Annual	Ecosystem Diagnostics IC-ECON	D1
▶ CDE-003 **Theory-of-Change fit** Alignment of the concept to the Theory of Change.	Categorical — code 397,000 rows · 99.0% populated · 4 distinct	High	Internal	**Validity** Assessed	ToC Knowledge Base IC-POLICY	D1 D3
▶ CDE-004 **Diagnostic data** Ecosystem diagnostic indicators.	Numeric — ratio 216,000 rows · 98.7% populated · 94,000 distinct	High	Internal	**Accuracy, Completeness** ≥ 95%	Ecosystem Diagnostics IC-ECON	D2
▶ CDE-005 **Market failures** Documented market failures.	Text — free text 935,000 rows · 98.4% populated · 579,700 distinct	Medium	Internal	**Accuracy** Evidenced	Research Repository IC-ECON	D2
▶ CDE-006 **Investment recommendation** The grant recommendation put forward.	Categorical — code 754,000 rows · 98.1% populated · 3 distinct	High	Confidential	**Completeness, Validity** 100%	Grant Management System IC-GRANT	D3 D5
▶ CDE-007 **Available capital** Endowment capital available to allocate.	Numeric — currency 573,000 rows · 97.8% populated · 427,000 distinct	High	Confidential	**Accuracy** 100%	Endowment Ledger	D3
▶ CDE-008 **Grantee governance**	Categorical — code	High	Confidential	**Completeness**	Due-Diligence File	D4

The Critical Data Elements (CDEs) each decision depends on — the specific governed data whose quality must be assured because a decision or AI agent relies on it — with their definitions, golden source, criticality tier, quality dimensions and the profiles, rules and contracts that keep them fit.

> **ARCHITECTURE MESSAGE**
>
> The data-quality foundation. Presents the specific critical data elements decisions depend on — definition, golden source, criticality tier, DQ profile and contract rules — the layer that determines whether an AI-supported decision can actually be trusted.

WHAT YOU'LL SEE

- A CDE catalogue with definition, data type, classification and criticality tier
- Which decisions each CDE feeds and its governing policy
- DQ profiles (the statistical shape of the data) and DQ rules carried in the data contract

HOW TO USE IT

- Open a decision to see the CDEs it relies on
- Read a CDE's profile to understand its real distribution
- Check the DQ rules and thresholds that govern it in the contract

Open Critical Data →

AI Agents

AI Agents & Models `ai_agents.html`

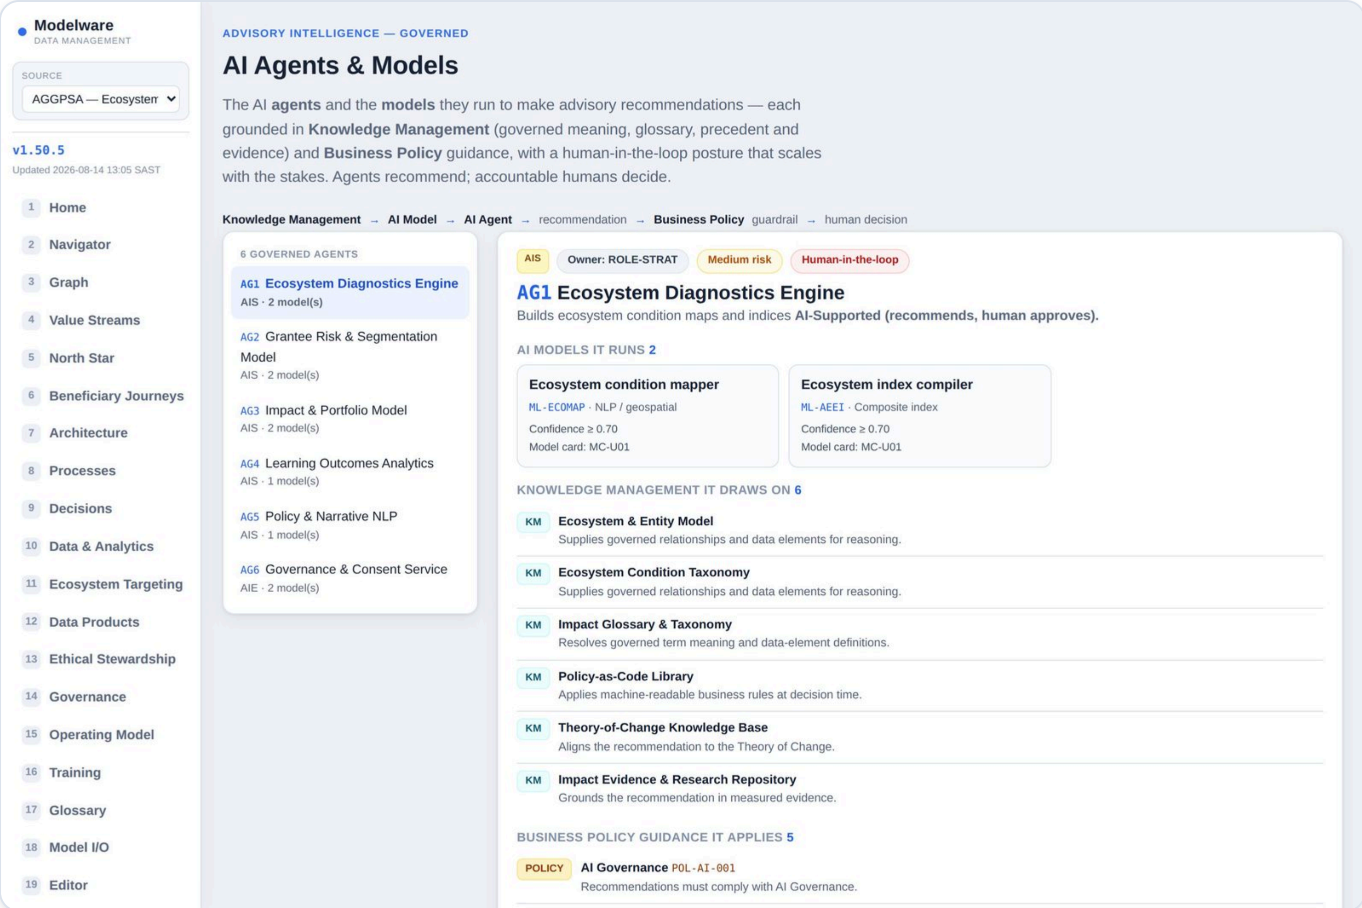

The AI agents and models that advise decisions and execute steps — each agent with the use-cases it serves, the AI models behind it, the knowledge assets it draws on, the policies that bind it, and its human-in-the-loop posture.

> **ARCHITECTURE MESSAGE**
>
> The AI-operations layer. Presents the running agents and the models behind them, with the knowledge assets, policies and human-in-the-loop posture that govern them — where accountability for live AI sits.

WHAT YOU'LL SEE

- An agent registry with the models, knowledge sources and policies for each
- The decisions each agent advises and the steps it executes
- Model cards, confidence thresholds, risk tier and HITL posture

HOW TO USE IT

- Open an agent to see what it advises and how it is governed
- Trace a model back to the knowledge and policies that constrain it
- Confirm the human-in-the-loop control on any advised decision

Open AI Agents →

Policy Inspector

Policy Inspector `policy_inspector.html`

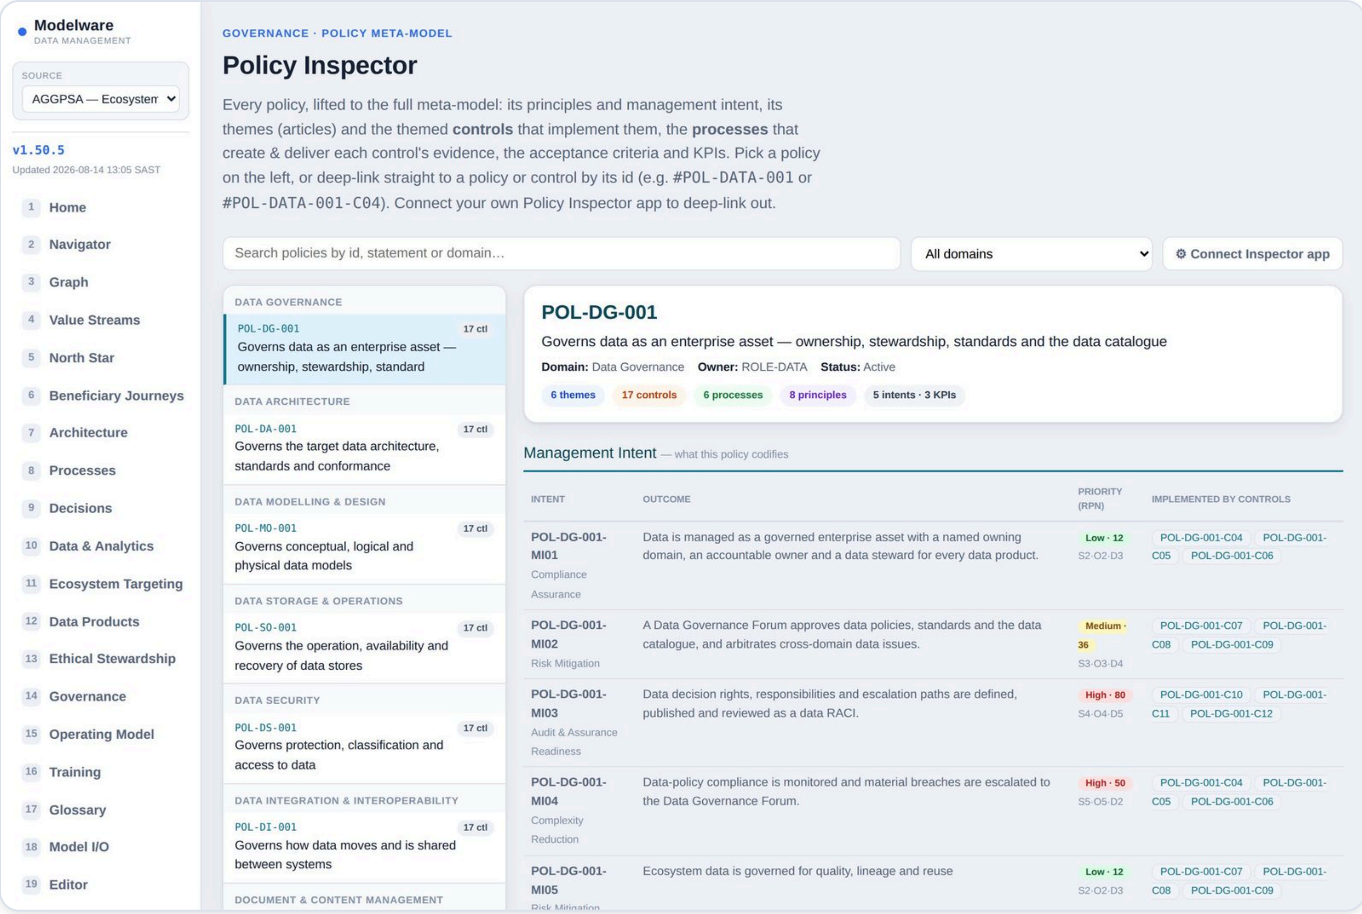

Inspect every policy in the library as a full meta-model — its Management Intent, Principles, Themes, Controls, minimum & assurance Evidence and RACI, in one place — and connect the controls out to your external Policy Inspector app. Filter to a domain to see only that policy's controls.

> **ARCHITECTURE MESSAGE**
>
> The policy layer, in full. Presents each policy as its governed meta-model — Management Intent → Principles → Themes → Controls → Evidence → RACI — the codified rules that turn governance intent into assured, evidenced, owned controls.

WHAT YOU'LL SEE

- A policy list (filter by domain or search) and a detail panel
- The visual hierarchy Policy → Theme → Control → Evidence → RACI
- Each control's Responsible (Business Data Steward + discipline Technical Specialist), Accountable owner, Consulted and Informed roles

HOW TO USE IT

- Pick a domain to narrow to its single authoritative policy
- Read a policy from intent down through its themed controls to the evidence that proves them
- Follow a control's link out to the external Policy Inspector app

Open Policy Inspector →

DM Maturity

Data Management Maturity `data_maturity.html`

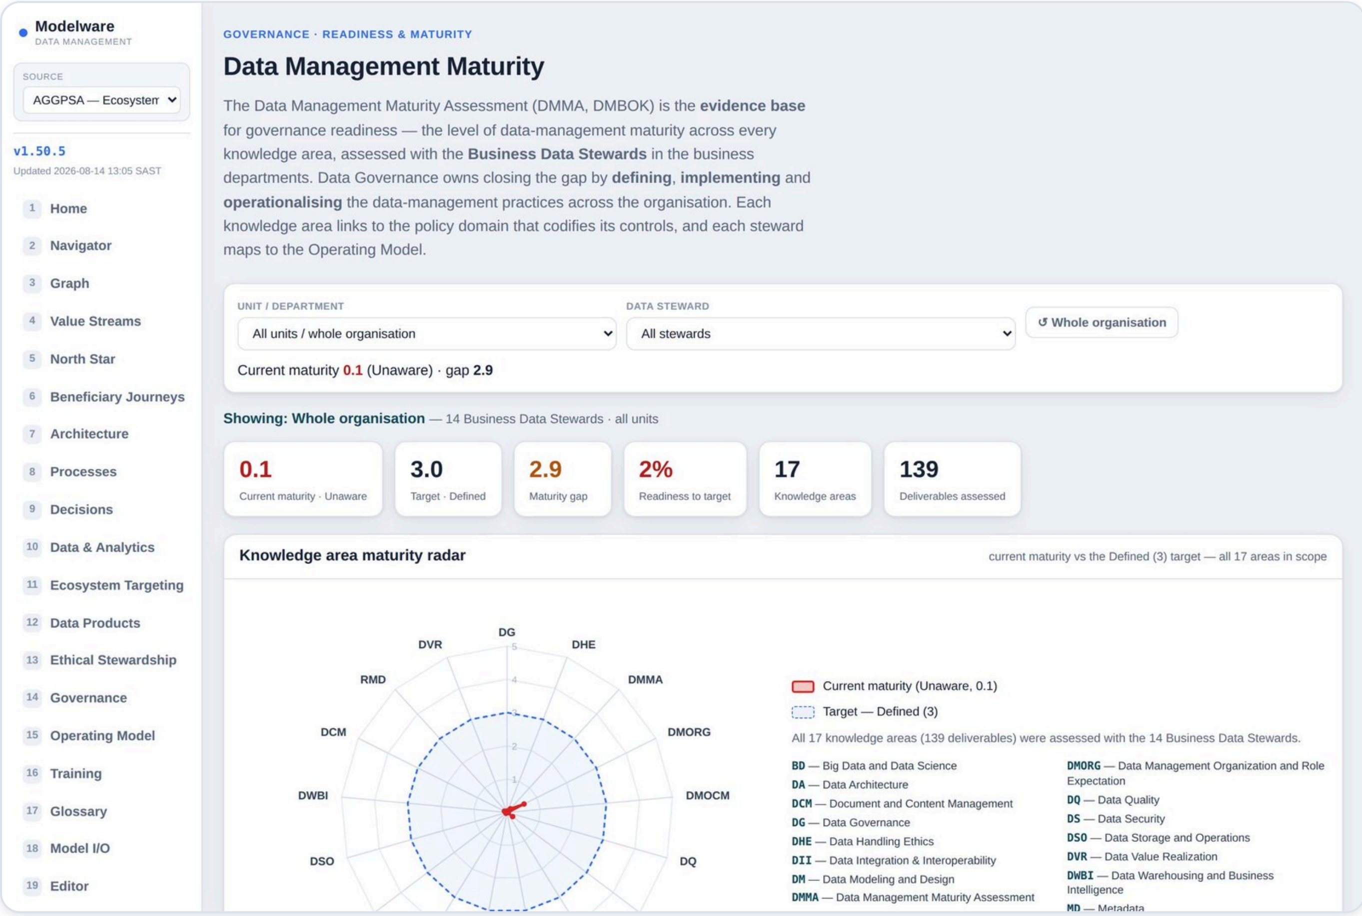

The Data Management Maturity Assessment (DMMA, DMBOK) — the governance-readiness evidence base. Current-vs-target maturity across every knowledge area, the maturity radar, the Define → Implement → Operationalise mandate, the Business Data Steward network, and the stewards' comments. Select a steward or a unit to see how they scored.

> **ARCHITECTURE MESSAGE**
>
> The governance-readiness evidence layer. Presents the measured DMMA maturity across every DMBOK knowledge area against target, mapped to the policy domains that codify each area's controls and to the Business Data Stewards who hold them — the baseline the whole governance programme is measured from and steers toward.

WHAT YOU'LL SEE

- A maturity radar of the 17 knowledge areas (current vs the Defined target), plus per-DMBOK-layer and per-area breakdowns mapped to their policy domains
- A Unit / Steward selector — pick one to recompute the whole view to that perspective, overlaid on the organisation average
- The Business Data Steward network by business unit and the 48 steward comments, filterable

HOW TO USE IT

- Read the overall gap first — current maturity vs the Defined (3) target
- Select a steward or unit to see how they scored, versus the org average
- Filter the comments to a domain or steward to hear the challenges in their words

Open DM Maturity →

Model Review

Architecture Model — Full Review `model_review.html`

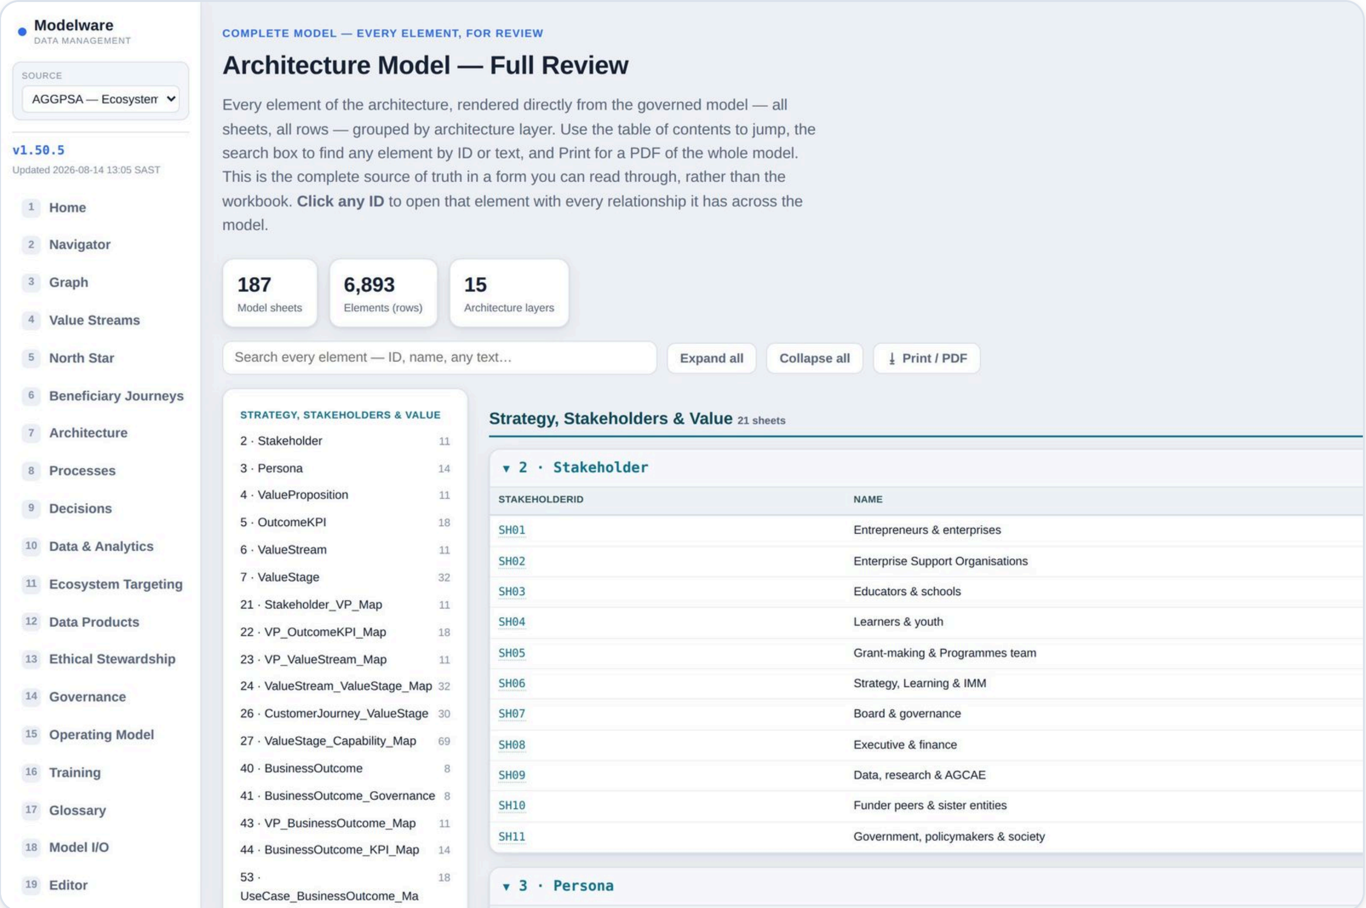

The whole architecture on one reviewable page — every sheet of the model laid out for review, grouped into 14 architecture layers, with search, a table of contents and a Print button, so the model can be reviewed completely without opening the Excel.

> **ARCHITECTURE MESSAGE**
>
> The whole model, made reviewable. Presents every element of every layer on one page, grouped into the fourteen architecture layers, so the complete architecture can be read and signed off without opening the spreadsheet.

WHAT YOU'LL SEE

- Every element of every sheet, grouped by architecture layer
- A searchable table of contents
- Clickable IDs that open an element's full detail panel

HOW TO USE IT

- Scan a layer, or search for an element or ID
- Click any ID to open its attributes and relationships
- Print the whole review to a PDF for sign-off

Open Model Review →

Relationships

Element Relationships `element_relationships.html`

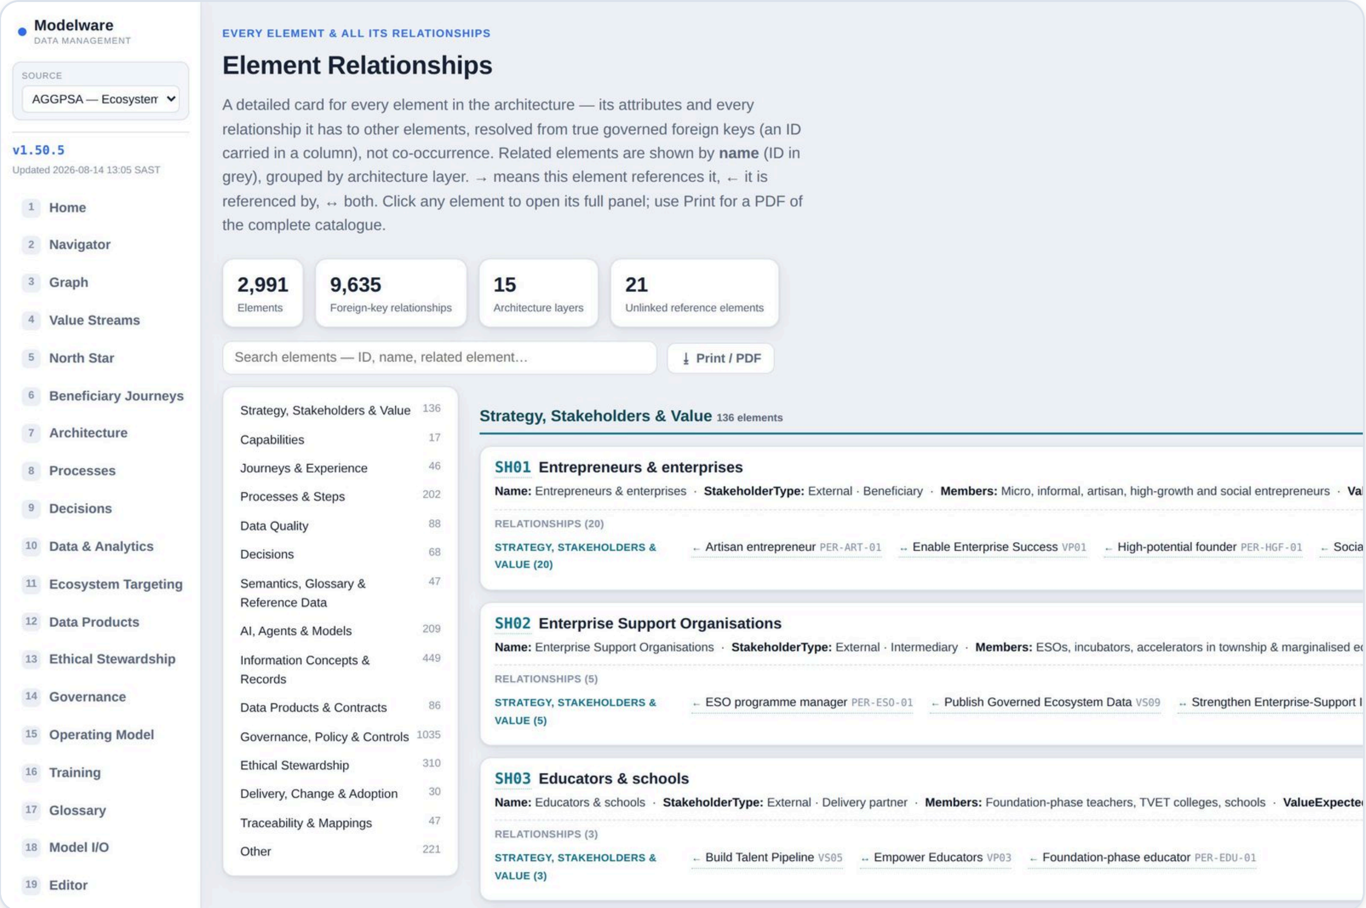

A detailed card for every element and all of its relationships — resolved from true governed foreign keys, shown by name with direction (references / referenced by / both) and grouped by architecture layer — so every element can be reviewed with its full context.

> **ARCHITECTURE MESSAGE**
>
> The connective tissue, made explicit. Presents every element with all of its true foreign-key relationships — by name and direction — the proof that this is one integrated model rather than a set of disconnected diagrams.

WHAT YOU'LL SEE

- A card per element with its attributes and every related element
- Relationships by name (ID in grey) with direction, grouped by layer
- A count of genuinely unlinked reference elements

HOW TO USE IT

- Search for an element and read its relationships
- Click any related element to traverse the model
- Print the full relationship catalogue to a PDF

Open Relationships →

Standalone artifact

Policy Meta-Model `policy_metamodel.html`

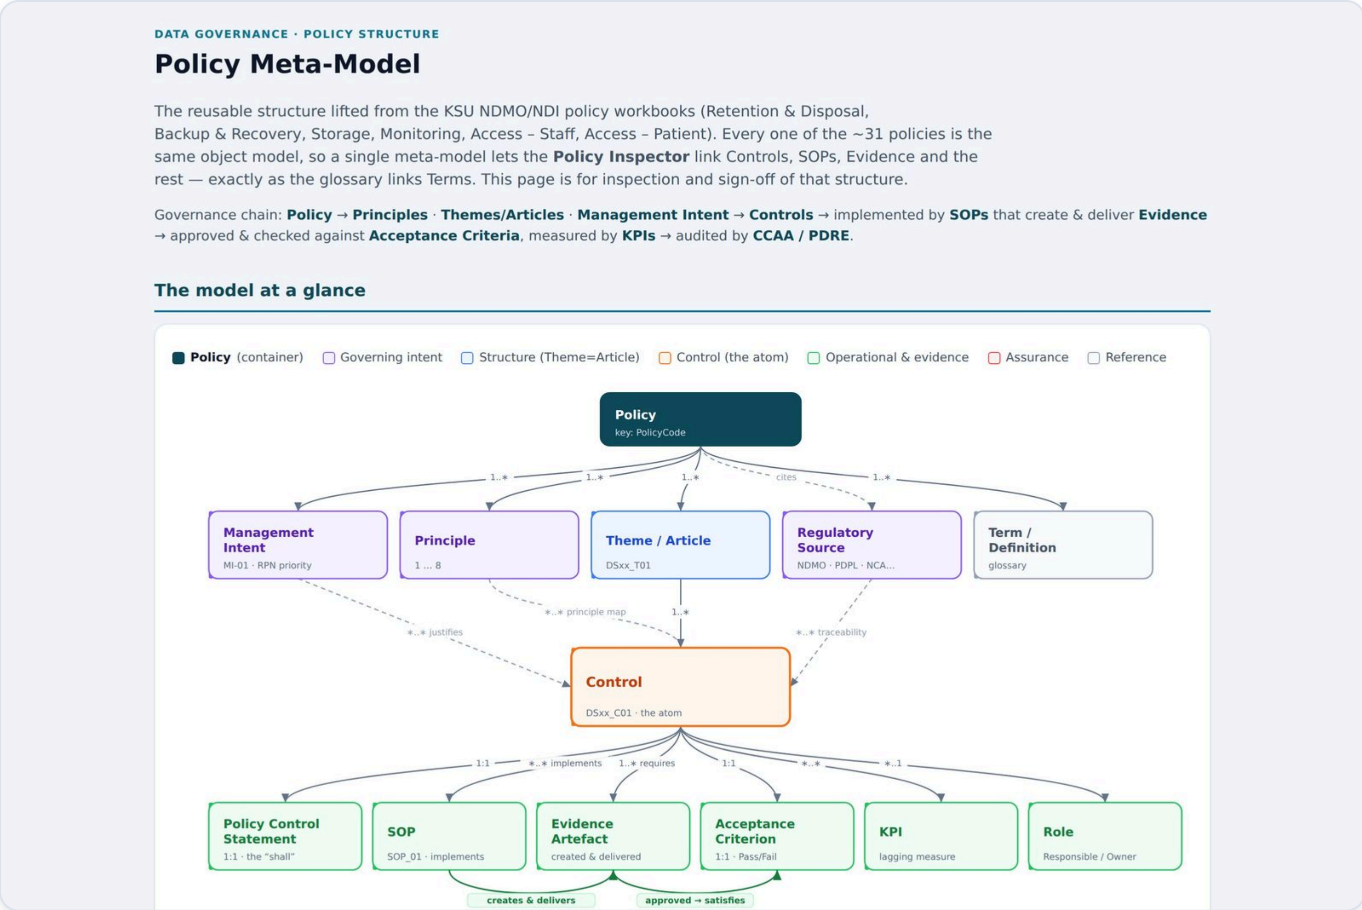

A standalone, self-contained view of the Policy meta-model — the same Policy → Theme → Control → Evidence → RACI structure as the Policy Inspector, packaged as a shareable artifact outside the pack navigation (delivered separately, not on the sidebar).

> **ARCHITECTURE MESSAGE**
>
> The policy layer as an inspectable meta-model — how a single policy is decomposed into the themed controls and the evidence that make governance intent auditable and assurable.

7 · Changing the model (for editors)

Two tabs change the model. Most people never need them — but for the approved few, this is how an edit reaches everyone.

MODEL EDITOR (APPROVED EDITORS)

- Passphrase-gated; read-only until unlocked.
- Relationship (…_Map) sheets are grids — add a row to add a relationship.
- Every other sheet is a form; add and delete rows anywhere.
- You edit a local draft, then *Save for repo* and *Export .xlsx*.

MODEL I/O (EVERYONE)

- Export the whole model to Excel to verify it.
- Re-import an edited workbook as a live source.
- A referential-integrity check flags any broken links.

> **Publish workflow:** pull the private repository → make the edit in the Editor → Save for repo (updates the model bundle) → commit → push. Because every page derives from that bundle, the change appears across the whole pack on the next load — no separate rebuild step.

8 · Element ID quick reference

IDs you'll see on chips and in the breadcrumb, and what they mean.

PREFIX	ELEMENT	EXAMPLE
`SH`	Stakeholder	`SH01`
`PER`	Persona	`PER03`
`VP`	Value Proposition	`VP01`
`BO / Outcome`	Business Outcome	`O1`
`K / KPI`	KPI	`CAC`
`VST`	Value Stream	`VST01`
`VSG / Stage`	Value Stage	`VSG04`
`CJ`	Customer Journey	`CJ-RET-01`
`C`	Capability	`C4`
`P`	Business Process	`P1`
`P…-S…`	Process Step	`P1-S01`
`D`	Decision	`D1`
`U`	AI Use-Case	`U01`
`AG`	AI Agent	`AG02`
`ML / AM`	AI Model	`ML-IMPACT`
`KA / SM`	Knowledge Asset / Semantic Model	`SM01`
`CDE`	Critical Data Element	`CDE-003`
`DP`	Data Product	`DP01`
`DM / Domain`	Data Domain	`DM03`
`IC`	Information Concept	`IC-CUST`
`RC`	Record Class	`RC-IMM`
`PD / POL`	Policy Domain / Policy	`PD-AI`
`CTL`	Control	`CTL-AI-001`
`EV`	Evidence	`EV-XAI-001`