

User Manual

How to read and use every part of this pack — the top-bar controls, the traceability spine, the Navigator flow, and each of the tabs. Written so anyone in Marketing, from an executive to a data steward, can find what they need and understand what they are looking at.

1 · Orientation — how the pack is built

This pack is one connected model. Every element traces along a single chain, and the same data drives every page — so a fact you agree on one tab is the same fact everywhere else.

THE PARALLEL-MAPPED MODEL

- **Value Stream** and **Customer Journey** are separate, first-class domains.
- They are joined by mapping tables and cross-mapped to **Capabilities**.
- This follows BIZBOK / TOGAF practice — the promise, the experience and the delivery are modelled side by side.

ONE SOURCE OF TRUTH

- All pages read from the same model bundle.
- An approved editor's change flows to every page automatically.
- The pack is self-contained and works offline — no server needed.

New to the pack? Read this section, then open the Navigator and walk one chain end-to-end. That single walk explains the whole model faster than any page of text.

2 · The controls you'll use everywhere

TOP NAVIGATION BAR

- Every tab is one click away along the top.
- The bar scrolls sideways if the window is narrow.
- The active tab is highlighted.

TRACEABILITY BREADCRUMB

- The strip under the nav shows the full chain, every step visible.
- As you make a selection it fills in with your chosen elements.
- Click any step to jump to it in the Navigator.

SOURCE SELECTOR (TOP-RIGHT)

- Switches the data the pack shows.
- *Generic (baked-in)* is the published model.
- *Imported (Excel)* appears once you import a workbook on the Model I/O page.

CARRIED SELECTION

- A selection made in the Navigator or on a chip is remembered.
- It follows you to every other page and narrows what you see.
- Clear it from the Navigator or the Graph's path controls.

The `version` and last-updated stamp sits at the right of the breadcrumb. The authoritative version control is the Git repository — each release is a commit.

3 · The Navigator flow — 13 steps end to end

The Navigator walks these levels in order. Each step narrows to what connects to your previous choice, and your path is carried to every other tab.

1

STEP 1

Stakeholder

Who the value is for — the stakeholder or persona. Everything downstream traces back to a stakeholder.

2

STEP 2

Value Proposition

The promise made to that stakeholder. Every stakeholder has at least one value proposition.

3

STEP 3

Business Outcome

The measurable outcome the proposition targets — what 'good' looks like, expressed as business outcomes and their KPIs.

4

STEP 4

Value Stream

The end-to-end value stream that delivers the proposition — a first-class domain in the parallel-mapped model.

5

STEP 5

Value Stage

A stage within the value stream, each with entrance/exit criteria and a primary KPI that rolls up to the outcome.

6

STEP 6

Customer Journey

How the customer experiences the stream — the customer journey, mapped stage-for-stage to the value stages.

7

STEP 7

Capability

The business capabilities the stream and journey require — the abilities the bank must have to deliver.

8

STEP 8

Business Process

The business processes that operationalise a capability — shown as SIPOC with their steps.

9**STEP 9****Decision**

The decisions taken inside a process — DMN-style rules, some automated, some kept human-in-the-loop.

10**STEP 10****AI Use Case**

The AI use-cases that automate or assist a decision, scored for value and risk.

11**STEP 11****AI Agent**

The AI agents that run a use-case, each carrying its human-in-the-loop stance.

12**STEP 12****Data Product**

The governed data products an agent needs — owned, contracted, and served from a domain.

13**STEP 13****Data Domain**

The data domain that groups and governs the data products — the owner/steward home for the data.

4 · Every tab, explained

Home

Landing

The front door. Explains the engagement, shows the architecture spine, lists every page with who it is for, links the downloadable workbooks and shows the current published version.

WHAT YOU'LL SEE

- The one-line architecture spine (Stakeholder → ... → KPIs)
- A directory card for every page, tagged by audience
- Publish status, active source and version history

HOW TO USE IT

- Start here to get your bearings
- Click any directory card to open that page
- Check the version history to see what changed and when

[Open Home →](#)

Navigator

Architecture Navigator

A guided, step-by-step walk through the whole parallel-mapped model. You pick a starting point and move one link at a time along the traceability chain, carrying your selection with you.

WHAT YOU'LL SEE

- A stepper for the 13 levels of the chain
- At each step, the elements available given what you've already chosen
- A breadcrumb that fills in as you go

HOW TO USE IT

- Pick a Stakeholder (or any step) to begin
- Move forward and back along the chain; your selection is carried to every other page
- Use it to answer 'what connects to what' end-to-end

[Open Navigator →](#)

Graph

Navigation Graph

An interactive picture of the model. Two views: an expandable hierarchy tree, and a Value Flow (Sankey) that shows how value flows from stakeholders down to data products.

WHAT YOU'LL SEE

- Tree view: expand any node to drill down every level
- Value Flow view: Stakeholder → Value Proposition → Value Stream → Value Stage → Capability → AI Use-Case → Data Product
- A 'Weight by' selector (AI value score or KPI roll-up) and a start-from stakeholder picker

HOW TO USE IT

- Click a node to focus it and read its explanation and value
- Switch to Value Flow to see where value concentrates
- If you've made a selection elsewhere, the flow narrows to just that path

[Open Graph →](#)

Architecture

Business Architecture

The capability centrepiece. Shows the capability map under a chosen value proposition, with value-stream heat and an Attention/Maturity toggle, plus the full capability grid for sign-off.

WHAT YOU'LL SEE

- A capability map coloured by heat
- An Attention vs Maturity toggle
- A drill-in panel for any capability

HOW TO USE IT

- Choose a value proposition to frame the map
- Toggle between Attention and Maturity
- Click a capability to drill in and sign it off

[Open Architecture →](#)

Journeys

Customer Journey & CX

The journey-map explorer across the lifecycle backbone, per line of business, with the CLV governance view and the CXO operating model.

WHAT YOU'LL SEE

- Journey tabs per line of business (Retail, Commercial, Wealth, Trade)
- Each stage with its experience score, touchpoints, capabilities, processes and KPIs
- The mapped value stream and business outcomes

HOW TO USE IT

- Pick a line of business tab
- Read left-to-right along the journey stages
- Click a KPI or stage to see detail

[Open Journeys →](#)

Processes

Business Processes (SIPOC)

Every high-level process as a SIPOC — Suppliers, Inputs, Process, Outputs, Customers — with its capabilities and participants, traced back to stakeholders.

WHAT YOU'LL SEE

- A SIPOC card per process
- The capabilities each process realises
- Participants and the stakeholders served

HOW TO USE IT

- Scan the SIPOC to understand a process at a glance
- Follow the capability links up to the architecture
- Use it to confirm ownership and hand-offs

[Open Processes →](#)

Decisions

Decision Models

Every decision as a DMN-style decision table — inputs, when → then rules, the outcome and its owner — with where each one is used.

WHAT YOU'LL SEE

- A decision table per decision
- Input conditions and the resulting outcome
- The owner and the processes/use-cases that call it

HOW TO USE IT

- Read the rules top-to-bottom
- Check the owner for accountability
- Trace a decision to the AI use-case that automates it

[Open Decisions →](#)

Personalisation

Hyper-personalisation & CDP

The closed-loop personalisation engine and the CDP capability model — the data, models and governance that make real-time personalisation work.

WHAT YOU'LL SEE

- The personalisation closed loop
- The CDP service model
- The governance guardrails around it

HOW TO USE IT

- Follow the loop from signal to action to measurement
- See which CDP services each use-case needs
- Use it to frame the data and model dependencies

[Open Personalisation →](#)

Data Products

Data Products

The governed, owned data products that deliver the AI use-cases and CDP — organised by data domain, each with an owner, steward and data contract.

WHAT YOU'LL SEE

- Data products grouped by domain
- Owner and steward per product
- A data-contract panel (schema, classification, sources)

HOW TO USE IT

- Filter by domain
- Click a product to open its data contract
- Trace a product up to the use-cases it serves

[Open Data Products →](#)

AI Use-Cases

AI Use-Cases

A value-vs-risk portfolio of marketing AI use cases, with live 1–5 driver scoring and an Excel round-trip.

WHAT YOU'LL SEE

- A value-vs-risk bubble chart of every use-case
- Live driver sliders (growth, relevance, retention, cost; data, model, integration, regulatory)
- Export/import of the scoring model

HOW TO USE IT

- Read the portfolio: top-right = high value, watch the risk axis
- Adjust drivers to test scenarios
- Export to Excel, score offline, import back

[Open AI Use-Cases →](#)

Training

Training & Adoption

Role pathways across four learning phases, with Champions as the Responsible-AI engine and the adoption feedback loop.

WHAT YOU'LL SEE

- Learning pathways per role
- Four phases from awareness to mastery
- The Champions model and feedback loop

HOW TO USE IT

- Find your role's pathway
- Follow the four phases in order
- See how Champions carry Responsible-AI into the business

[Open Training →](#)

Value Streams

Value Streams — Ownership & Governance

Trace each value stream down to who owns and governs it — value stages and KPIs, the business outcomes it delivers with their accountable roles and forums, the ownership-role model, and the governance operating layer.

WHAT YOU'LL SEE

- A value-stream picker
- Value stages with primary KPIs
- Business outcomes with accountable roles and review forums
- Councils, decision-rights RACI, policy → control, and the risk register

HOW TO USE IT

- Pick a value stream on the left
- Read down from stages and KPIs to outcomes and owners
- Use the governance layer to confirm accountability

[Open Value Streams →](#)

Governance

Governance & Responsible AI

A layered heatmap of Data Governance, Responsible AI and Change & Adoption, with a success-measure scorecard.

WHAT YOU'LL SEE

- A layered governance heatmap
- Responsible-AI controls
- A success-measure scorecard

HOW TO USE IT

- Read the heat to spot weak layers
- Check the scorecard against targets
- Use it to brief governance forums

[Open Governance →](#)

Model I/O

Model Export / Import

Export the complete model to Excel — one sheet per level, every attribute and link, with a referential-integrity check — and import an edited workbook back as a live source that drives every page.

WHAT YOU'LL SEE

- An Export to Excel action
- An Import from Excel action
- A referential-integrity validation summary

HOW TO USE IT

- Export to verify the model in a spreadsheet
- Edit and re-import to load it as the live 'Imported (Excel)' source
- Use the integrity check to catch broken links

[Open Model I/O →](#)

Editor

Model Editor

Edit every element of the model — all sheets. Relationship (..._Map) sheets are editable grids where you add a relationship; every other sheet is a form. Restricted to approved editors, and changes are published through the private repository.

WHAT YOU'LL SEE

- A passphrase gate (approved editors only)
- Grids for relationship sheets; forms for everything else
- Add / delete rows anywhere; Save for repo, Export .xlsx, Load published

HOW TO USE IT

- Unlock with the editor passphrase
- Edit a local draft in the browser
- Save for repo, then pull → commit → push to publish to everyone

[Open Editor →](#)

Glossary

Glossary settings

Point the pack at the Nedbank Glossary Workbench — base URL and link template — so governed terms deep-link out to their definitions.

WHAT YOU'LL SEE

- A base-URL setting
- A link-template setting
- A live preview of a resolved term link

HOW TO USE IT

- Set the glossary base URL and template
- Save the settings
- Governed terms across the pack then deep-link to the Workbench

[Open Glossary →](#)

5 · Editing the model and publishing changes

Two tabs change the model. Most people never need them — but for the approved few, this is how an edit reaches everyone.

MODEL EDITOR (APPROVED EDITORS)

- Passphrase-gated; read-only until unlocked.
- Relationship (..._Map) sheets are grids — add a row to add a relationship.
- Every other sheet is a form; add and delete rows anywhere.
- You edit a local draft, then *Save for repo* and *Export .xlsx*.

MODEL I/O (EVERYONE)

- Export the whole model to Excel to verify it.
- Re-import an edited workbook as a live source.
- A referential-integrity check flags any broken links.

Publish workflow: pull the private repository → make the edit in the Editor → Save for repo (updates the model bundle) → commit → push. Because every page derives from that bundle, the change appears across the whole pack on the next load — no separate rebuild step.

6 · Element ID quick reference

IDs you'll see on chips and in the breadcrumb, and what they mean.

PREFIX	ELEMENT	EXAMPLE
SH	Stakeholder	SH01
PER	Persona	PER03
VP	Value Proposition	VP01
BO / Outcome	Business Outcome	01
K / KPI	KPI	CAC
VST	Value Stream	VST01
VSG / Stage	Value Stage	VSG04
CJ	Customer Journey	CJ-RET-01
C	Capability	C4
P	Business Process	P1
D	Decision	D1
U	AI Use-Case	U01
AG	AI Agent	AG02
DP	Data Product	DP01
DM / Domain	Data Domain	DM03
SM	Semantic Model	SM01