

User Guide & primer

A plain-language guide to using this pack **and** to the ideas behind it. It assumes no prior knowledge of business architecture. If the words *value stream*, *capability* or *data product* are new to you, start at “The big idea” and read straight down — by the end you'll be able to look at any AI initiative and say exactly how it connects to a business outcome, what it depends on, and how you'd know it worked. That is the whole point of the pack: to make Data & AI decisions legible to the people who sponsor them.

1 · The big idea — why AI needs a business architecture

Most AI initiatives fail for a business reason, not a technical one.

They get built without a clear line back to a decision that matters, an outcome someone owns, or data anyone trusts. A **business architecture** is simply the map that draws those lines. It lets you point at any AI idea and ask: *whom does it serve, what decision does it change, what result does it move, and what does it depend on?* If the map has a gap, that is where the initiative will struggle.

This pack lays out one such map for Nedbank Marketing. Everything in it hangs off a single chain we call the **golden thread** — a line from a person the bank serves, all the way down to the governed data an AI system runs on:



WHAT YOU CAN DO WITH THE THREAD

- **Justify** an AI initiative: trace it up to the outcome and KPI it moves.
- **De-risk** it: trace it down to the data it needs and the human oversight it keeps.
- **Prioritise**: compare initiatives on the value they create versus the risk they carry.
- **Align** a room: give everyone the same words for the same things.

HOW TO READ THIS GUIDE

- Section 2 defines every building block in plain language, with a running example.
- Section 3 is a short playbook for taking one AI idea through the model.
- Then: how to read the charts, how to move around, and what each tab is for.
- No jargon is assumed — terms are defined the first time they appear.

Prefer to learn by doing? Open the Navigator and walk one thread end-to-end. Pick a stakeholder and keep clicking forward — the guide below explains what each step means as you go.

2 · The building blocks, in plain language

These are the thirteen links of the golden thread, in order. For each one: what it is, the question it answers, why it matters to a sponsor, and a worked example from a single running story — **preventing customer churn**. Read them top to bottom and the whole model falls into place. Wherever AI enters, it is called out in amber.

1 Stakeholder Who is this for?

What it is. The people or groups the business serves or must answer to — customers and prospects, partners, regulators, and internal teams like Marketing or Risk.

Why it matters. An AI initiative that can't name the stakeholder it helps is a solution looking for a problem. Everything in the model traces back to a stakeholder, which is why the walk starts here.

Example. *Customers & prospects — the people the bank wants to acquire, serve well, and keep.*

2 Value Proposition What do we promise them?

What it is. The value proposition — the promise of value made to that stakeholder. What they get, and why it is worth their attention.

Why it matters. It anchors every downstream cost and project to a benefit someone actually wants. No promise, no reason to build.

Example. *“Relevant, trusted financial solutions at the right moment, through the right channel.”*

3 Business Outcome What result are we trying to move — and how will we know?

What it is. The business outcome is the measurable result that proves the promise is being kept. Each outcome carries KPIs — the specific numbers that quantify it.

Why it matters. This is where AI has to earn its place. If an initiative can't name the outcome and KPI it moves, it can't be justified or measured. Outcomes and KPIs are the language executives already speak.

Example. *Outcome: increase customer lifetime value. KPI: retention rate, and lifetime value per customer.*

4 Value Stream How does the business create this value, start to finish?

What it is. A value stream is the end-to-end set of activities that creates and delivers the value — from the trigger to the result — seen from the organisation's side.

Why it matters. It gives an AI investment somewhere concrete to live. You don't improve "the business" in the abstract; you improve a stream.

Example. *Develop Customer Relationship — everything from first contact to a deep, retained relationship.*

5

Value Stage

Which specific step are we improving?

What it is. A value stage is one step of the stream, with a clear start and finish and its own KPI. Stages are how a long stream is broken into measurable pieces.

Why it matters. AI is applied stage by stage. Because each stage's KPI rolls up into the outcome, you can see exactly where value is won or lost — and target the weakest step.

Example. *Retain Relationship — the stage where churn is prevented. KPI: retention rate.*

6

Customer Journey

What does this feel like for the customer?

What it is. The customer journey is the same value seen from the customer's side — the experience they actually go through, moment by moment.

Why it matters. AI that improves an internal number but worsens the experience destroys value. Modelling the journey alongside the stream keeps the customer in the picture.

Example. *A customer quietly considers leaving; the bank notices and reaches out with a timely, relevant offer — rather than a generic mailshot.*

7

Capability

What must we be able to do?

What it is. A capability is an ability the business must have to deliver — a "what we can do," independent of how teams are organised or which tools they use.

Why it matters. Capabilities are reusable building blocks. One AI investment in a capability can serve many streams and journeys, which is how you avoid re-solving the same problem five times.

Example. *CLV Management and Customer Intelligence — the abilities to understand, predict and grow customer value.*

8

Business Process

How exactly do we do it today?

What it is. A business process is the concrete, repeatable way a capability actually gets done — its steps, inputs and outputs. The pack shows each as a SIPOC (Suppliers, Inputs, Process, Outputs, Customers).

Why it matters. You cannot automate what you cannot describe. The process is the real work where AI plugs in — and where you see what a model would change.

Example. *CLV Optimisation — the process that scores customers for risk and value, then triggers the right retention action.*

9 Decision

What is being decided, and on what basis?

What it is. A decision is a single judgement made inside a process: inputs go in, an outcome comes out, following rules. The pack shows each as a decision table (if these conditions, then this outcome).

Why it matters. Most AI automates or assists a decision. Naming the decision is the exact moment a vague “AI initiative” becomes something concrete you can build, test and govern.

Example. *Retention Decision — is this customer at risk, and what is the best save action for them?*

Where AI comes in. This is the hand-off point to AI: a rule-based decision becomes a candidate for a model that decides faster, at scale, and more accurately.

10 AI Use Case

What is the AI actually doing for us?

What it is. An AI use-case is a specific application of AI to a decision or task — with a value it creates and a risk it carries. This is the “AI initiative” itself.

Why it matters. By this point the initiative is anchored to a decision, a stage, an outcome and a KPI — so its value and its risk are both visible and comparable against every other use-case, instead of being sold on hype.

Example. *Churn-risk prediction & save — predict who is likely to leave and recommend the best retention offer.*

Where AI comes in. In this pack each use-case is scored 1–5 on value (growth, relevance, retention, CX, cost) and on risk (data readiness, model, integration, regulatory) — so a portfolio can be prioritised on evidence.

11 AI Agent

What runs it, and who oversees it?

What it is. An AI agent is the working system that runs a use-case — the model or service that does the job — together with its stance on human oversight.

Why it matters. Executives need to know not just that AI is used, but how it is operated and controlled. The agent is where accountability for the running system lives.

Example. *Propensity & CLV Scoring agent — scores customers daily; high-value retention offers are routed to a banker for approval.*

Where AI comes in. Human-in-the-Loop: for high-stakes or high-risk outputs, a person reviews and approves before anything acts. It is the main dial between full automation and human control.

12 Data Product

What data does this depend on — and can we trust it?

What it is. A data product is a governed, owned, reusable data asset the AI needs, with a clear owner, an agreed definition, and a quality/contract. Not a raw table — a managed product.

Why it matters. AI is only as good as the data beneath it. Data products make that dependency explicit and accountable — often the real difference between a demo that impresses and a system that survives in production.

Example. *Customer Identity Graph and Model Score Products — the trusted data the churn model consumes and the scores it produces.*

13 Data Domain

Who is accountable for this data?

What it is. A data domain is the area of the business that owns and governs a group of related data products — with a named owner and steward.

Why it matters. It puts a name against the data. Ownership and stewardship are what make data trustworthy at scale, and what a regulator or auditor will ask for first.

Example. *Insight & Model — the domain that owns customer scores and model outputs across Marketing.*

A few supporting terms

Words you'll meet across the pack that aren't their own step in the walk.

PERSONA

A named, representative example of a stakeholder — e.g. a specific customer segment. Personas make an abstract stakeholder concrete so teams design for a real person, not an average.

CDP (CUSTOMER DATA PLATFORM)

The engine that unifies customer signals in near-real-time and activates them — the plumbing that lets a “next best action” fire at the right moment. In the model it is a set of capabilities and services the personalisation use-cases draw on.

SEMANTIC MODEL / GOVERNED TERM

The agreed, written definition of a business term — what “churn”, “active customer” or “lifetime value” actually mean. When everyone's AI and every report uses the same definition, the numbers reconcile. This is the vocabulary layer beneath the data products.

HUMAN-IN-THE-LOOP (HITL)

A control where a person reviews or approves an AI output before it takes effect. Used for high-value, high-risk or regulated decisions. The presence or absence of HITL is one of the clearest signals of how much trust you are placing in a model.

The one sentence to remember: a good AI initiative is a *use-case* that automates a *decision* inside a *process*, which delivers a *capability* a *value stage* needs, moving a *KPI* for an *outcome* promised to a *stakeholder* — running on a governed *data product*, with a human in the loop where the stakes are high.

3 · How to apply AI to a business problem

A simple, repeatable way to take one idea through the model. Do it on paper, or do it live in the Navigator. The example in blue continues the churn story.

1 Start from an outcome, not a technology

Name the business result you want to move and the KPI that measures it. Resist starting with “let’s use GenAI.” Start with the number.

Outcome: grow customer lifetime value · **KPI:** retention rate.

2 Find the decision that moves it

Trace to the value stage that most affects that KPI, then to the decision made there. AI acts on decisions — so this is where a real initiative forms.

Stage: Retain Relationship → **Decision:** who is at risk and what is the best save?

3 Name the AI use-case

State plainly what the AI will do for that decision, and score it: how much value (growth, relevance, retention, CX, cost) against how much risk (data, model, integration, regulatory).

Use-case: churn-risk prediction & save · high value, moderate risk.

4 Check the data it depends on

List the data products the use-case needs and whether they are owned, defined and trustworthy. This is where most initiatives quietly fail — surface it early.

Needs: Customer Identity Graph, model scores · **owner:** Insight & Model domain.

5 Decide the human oversight

Choose where a person stays in the loop. High-value or regulated actions get human approval; low-stakes, high-volume ones can run automatically.

Auto-score everyone; a banker approves offers above a value threshold.

6 Close the loop on the KPI

Ship it, then watch the same KPI you started with. If it moves, you have proof; if not, you have a precise place to look. The thread makes the feedback loop obvious.

Track retention on the treated group vs a control; feed results back into targeting.

Notice what the playbook prevents: an AI project with no owned outcome, no named decision, or no trustworthy data. Those three gaps account for most failed initiatives — and the model surfaces all three before a line of code is written.

4 · Reading the visualisations

The pack shows the same model several ways. Here is how to read each one.

NAVIGATOR — THE GUIDED WALK

- Pick an element at each step; the next step shows only what connects to it.
- The numbered stepper on the left is your position along the thread.
- Best for learning the model and answering “what connects to what.”

GRAPH — VALUE FLOW (SANKEY)

- Ribbons flow left → right from Stakeholder down to Data Product.
- **Ribbon width = how much value flows** (AI value score, or KPI weight — your choice).
- Wider bands are where value concentrates. Use the filter rail to hold “all” or a few items at any level.

ARCHITECTURE — CAPABILITY HEATMAP

- Tiles are capabilities; colour shows attention needed or maturity.
- Hot tiles are where investment (including AI) will pay back most.

GRAPH — HIERARCHY TREE

- Expand any node (►) to drill down a level at a time.
- Click a node for its plain-language explanation and its value.
- Best for exploring freely rather than following one path.

AI USE-CASES — VALUE VS RISK

- Each bubble is a use-case; up-and-right = high value.
- Watch the risk axis: a high-value, high-risk idea needs stronger controls.
- Best for prioritising a portfolio on evidence, not opinion.

VALUE STREAMS — GOVERNANCE VIEW

- Traces a stream down to who owns and governs it.
- Best for answering “who is accountable” for an outcome or control.

5 · Finding your way around

THE LEFT SIDEBAR

- Every page is a numbered tab down the left; the current one is highlighted.
- The Modelware brand, the data *Source* selector and the version stamp live in the rail.
- On a narrow screen the rail folds into a bar across the top.

THE TRACEABILITY PATH

- Shown on the **Navigator** — it is the thread you are building.
- It fills in as you choose an element at each step.
- Other pages stay clean; your selection is still remembered behind the scenes.

THE DATA SOURCE SELECTOR

- Switches which data the whole pack shows.
- *Generic (baked-in)* is the published model.
- *Imported (Excel)* appears after you import a workbook on Model I/O.

CARRIED SELECTION

- A choice made in the Navigator follows you to other pages and focuses them.
- The Graph's Value Flow starts from it, and you can widen from there.
- Clear it with the Navigator's reset or the Graph's “Whole model”.

6 · Every tab — where to go for what

Home

Landing

The front door. Explains the engagement, shows the architecture spine, lists every page with who it is for, links the downloadable workbooks and shows the current published version.

WHAT YOU'LL SEE

- The one-line architecture spine (Stakeholder → ... → KPIs)
- A directory card for every page, tagged by audience
- Publish status, active source and version history

HOW TO USE IT

- Start here to get your bearings
- Click any directory card to open that page
- Check the version history to see what changed and when

[Open Home →](#)

Navigator

Architecture Navigator

A guided, step-by-step walk through the whole parallel-mapped model. You pick a starting point and move one link at a time along the traceability chain, carrying your selection with you.

WHAT YOU'LL SEE

- A stepper for the 13 levels of the chain
- At each step, the elements available given what you've already chosen
- A breadcrumb that fills in as you go

HOW TO USE IT

- Pick a Stakeholder (or any step) to begin
- Move forward and back along the chain; your selection is carried to every other page
- Use it to answer 'what connects to what' end-to-end

[Open Navigator →](#)

Graph

Navigation Graph

An interactive picture of the model. Two views: an expandable hierarchy tree, and a Value Flow (Sankey) that shows how value flows from stakeholders down to data products.

WHAT YOU'LL SEE

- Tree view: expand any node to drill down every level
- Value Flow view: Stakeholder → Value Proposition → Value Stream → Value Stage → Capability → AI Use-Case → Data Product
- A 'Weight by' selector (AI value score or KPI roll-up) and a start-from stakeholder picker

HOW TO USE IT

- Click a node to focus it and read its explanation and value
- Switch to Value Flow to see where value concentrates
- If you've made a selection elsewhere, the flow narrows to just that path

[Open Graph →](#)

Architecture

Business Architecture

The capability centrepiece. Shows the capability map under a chosen value proposition, with value-stream heat and an Attention/Maturity toggle, plus the full capability grid for sign-off.

WHAT YOU'LL SEE

- A capability map coloured by heat
- An Attention vs Maturity toggle
- A drill-in panel for any capability

HOW TO USE IT

- Choose a value proposition to frame the map
- Toggle between Attention and Maturity
- Click a capability to drill in and sign it off

[Open Architecture →](#)

Journeys

Customer Journey & CX

The journey-map explorer across the lifecycle backbone, per line of business, with the CLV governance view and the CXO operating model.

WHAT YOU'LL SEE

- Journey tabs per line of business (Retail, Commercial, Wealth, Trade)
- Each stage with its experience score, touchpoints, capabilities, processes and KPIs
- The mapped value stream and business outcomes

HOW TO USE IT

- Pick a line of business tab
- Read left-to-right along the journey stages
- Click a KPI or stage to see detail

[Open Journeys →](#)

Processes

Business Processes (SIPOC)

Every high-level process as a SIPOC — Suppliers, Inputs, Process, Outputs, Customers — with its capabilities and participants, traced back to stakeholders.

WHAT YOU'LL SEE

- A SIPOC card per process
- The capabilities each process realises
- Participants and the stakeholders served

HOW TO USE IT

- Scan the SIPOC to understand a process at a glance
- Follow the capability links up to the architecture
- Use it to confirm ownership and hand-offs

[Open Processes →](#)

Decisions

Decision Models

Every decision as a DMN-style decision table — inputs, when → then rules, the outcome and its owner — with where each one is used.

WHAT YOU'LL SEE

- A decision table per decision
- Input conditions and the resulting outcome
- The owner and the processes/use-cases that call it

HOW TO USE IT

- Read the rules top-to-bottom
- Check the owner for accountability
- Trace a decision to the AI use-case that automates it

[Open Decisions →](#)

Personalisation

Hyper-personalisation & CDP

The closed-loop personalisation engine and the CDP capability model — the data, models and governance that make real-time personalisation work.

WHAT YOU'LL SEE

- The personalisation closed loop
- The CDP service model
- The governance guardrails around it

HOW TO USE IT

- Follow the loop from signal to action to measurement
- See which CDP services each use-case needs
- Use it to frame the data and model dependencies

[Open Personalisation →](#)

Data Products

Data Products

The governed, owned data products that deliver the AI use-cases and CDP — organised by data domain, each with an owner, steward and data contract.

WHAT YOU'LL SEE

- Data products grouped by domain
- Owner and steward per product
- A data-contract panel (schema, classification, sources)

HOW TO USE IT

- Filter by domain
- Click a product to open its data contract
- Trace a product up to the use-cases it serves

[Open Data Products →](#)

AI Use-Cases

AI Use-Cases

A value-vs-risk portfolio of marketing AI use cases, with live 1–5 driver scoring and an Excel round-trip.

WHAT YOU'LL SEE

- A value-vs-risk bubble chart of every use-case
- Live driver sliders (growth, relevance, retention, cost; data, model, integration, regulatory)
- Export/import of the scoring model

HOW TO USE IT

- Read the portfolio: top-right = high value, watch the risk axis
- Adjust drivers to test scenarios
- Export to Excel, score offline, import back

[Open AI Use-Cases →](#)

Training

Training & Adoption

Role pathways across four learning phases, with Champions as the Responsible-AI engine and the adoption feedback loop.

WHAT YOU'LL SEE

- Learning pathways per role
- Four phases from awareness to mastery
- The Champions model and feedback loop

HOW TO USE IT

- Find your role's pathway
- Follow the four phases in order
- See how Champions carry Responsible-AI into the business

[Open Training →](#)

Value Streams

Value Streams — Ownership & Governance

Trace each value stream down to who owns and governs it — value stages and KPIs, the business outcomes it delivers with their accountable roles and forums, the ownership-role model, and the governance operating layer.

WHAT YOU'LL SEE

- A value-stream picker
- Value stages with primary KPIs
- Business outcomes with accountable roles and review forums
- Councils, decision-rights RACI, policy → control, and the risk register

HOW TO USE IT

- Pick a value stream on the left
- Read down from stages and KPIs to outcomes and owners
- Use the governance layer to confirm accountability

[Open Value Streams →](#)

Governance

Governance & Responsible AI

A layered heatmap of Data Governance, Responsible AI and Change & Adoption, with a success-measure scorecard.

WHAT YOU'LL SEE

- A layered governance heatmap
- Responsible-AI controls
- A success-measure scorecard

HOW TO USE IT

- Read the heat to spot weak layers
- Check the scorecard against targets
- Use it to brief governance forums

[Open Governance →](#)

Model I/O

Model Export / Import

Export the complete model to Excel — one sheet per level, every attribute and link, with a referential-integrity check — and import an edited workbook back as a live source that drives every page.

WHAT YOU'LL SEE

- An Export to Excel action
- An Import from Excel action
- A referential-integrity validation summary

HOW TO USE IT

- Export to verify the model in a spreadsheet
- Edit and re-import to load it as the live 'Imported (Excel)' source
- Use the integrity check to catch broken links

[Open Model I/O →](#)

Editor

Model Editor

Edit every element of the model — all sheets. Relationship (..._Map) sheets are editable grids where you add a relationship; every other sheet is a form. Restricted to approved editors, and changes are published through the private repository.

WHAT YOU'LL SEE

- A passphrase gate (approved editors only)
- Grids for relationship sheets; forms for everything else
- Add / delete rows anywhere; Save for repo, Export .xlsx, Load published

HOW TO USE IT

- Unlock with the editor passphrase
- Edit a local draft in the browser
- Save for repo, then pull → commit → push to publish to everyone

[Open Editor →](#)

Glossary

Glossary settings

Point the pack at the Nedbank Glossary Workbench — base URL and link template — so governed terms deep-link out to their definitions.

WHAT YOU'LL SEE

- A base-URL setting
- A link-template setting
- A live preview of a resolved term link

HOW TO USE IT

- Set the glossary base URL and template
- Save the settings
- Governed terms across the pack then deep-link to the Workbench

[Open Glossary →](#)

7 · Changing the model (for editors)

Two tabs change the model. Most people never need them — but for the approved few, this is how an edit reaches everyone.

MODEL EDITOR (APPROVED EDITORS)

- Passphrase-gated; read-only until unlocked.
- Relationship (..._Map) sheets are grids — add a row to add a relationship.
- Every other sheet is a form; add and delete rows anywhere.
- You edit a local draft, then *Save for repo* and *Export .xlsx*.

MODEL I/O (EVERYONE)

- Export the whole model to Excel to verify it.
- Re-import an edited workbook as a live source.
- A referential-integrity check flags any broken links.

Publish workflow: pull the private repository → make the edit in the Editor → Save for repo (updates the model bundle) → commit → push. Because every page derives from that bundle, the change appears across the whole pack on the next load — no separate rebuild step.

8 · Element ID quick reference

IDs you'll see on chips and in the breadcrumb, and what they mean.

PREFIX	ELEMENT	EXAMPLE
SH	Stakeholder	SH01
PER	Persona	PER03
VP	Value Proposition	VP01
BO / Outcome	Business Outcome	01
K / KPI	KPI	CAC
VST	Value Stream	VST01
VSG / Stage	Value Stage	VSG04
CJ	Customer Journey	CJ-RET-01
C	Capability	C4
P	Business Process	P1
D	Decision	D1
U	AI Use-Case	U01
AG	AI Agent	AG02
DP	Data Product	DP01
DM / Domain	Data Domain	DM03
SM	Semantic Model	SM01